

Notes: Abaluck & Adams-Prassl (QJE, 2021)
What Do Consumers Consider Before They Choose? Identification from
Asymmetric Demand Responses

Contents

1	Big picture	3
2	Incremental contribution of the paper	3
3	Data and measurement	4
3.1	Lab experiment	4
3.2	Medicare Part D application	4
3.3	Plan attributes	4
3.4	Key attention objects	5
4	Models and empirical specifications	5
4.1	General consideration-set framework	5
4.2	Daly-Zachary / full-consideration symmetry	5
4.3	Identification from asymmetric responses	5
4.4	Alternative-specific consideration model	6
4.5	Hybrid Medicare Part D model	6
4.6	Smart-default welfare simulations	6
5	Identification and empirical design	6
5.1	Core identification problem	6
5.2	Cross-price asymmetry solution	6
5.3	Why this is powerful	7
5.4	Validation and application logic	7
5.5	Limitations	7
6	Section-by-section study guide	7
7	Tables and figures	7
7.1	Figure I: Lab Experiment Sample Product Selection Screen + Figure II: Product Fixed Effects in Attention	7
7.2	Table I: Experimental Data Estimation Results	8
7.3	Table II: Sample Demographics and Plan Characteristics + Table III: Excess Sensitivity to Default Attributes in Switching Model	9
7.4	Table IV: Conditional Logit and Hybrid Model	10
7.5	Figure III: Distribution of Savings + Table V: Welfare Impact of a Smart Default Policy . . .	11
7.6	Table VI: Welfare Effect of a Smart Default Policy: Restricted Reassignment to Those with Savings Greater than Switching Costs	12

8	Result map	13
9	Short summary	13
10	Conclusion	14
10.1	Core conclusion	14
10.2	Economic intuition	14
10.3	Incremental contribution revisited	14
10.4	Policy implications	14
10.5	Limitations	14
10.6	Takeaway	14

1 Big picture

- **Question.** Can researchers identify what consumers consider before they choose when consideration sets are unobserved?
- **Core challenge.** Standard discrete-choice models assume that consumers consider all available alternatives. If that assumption fails, estimated preferences can confound utility with attention. A product may be chosen often because it is highly valued or because it is disproportionately considered.
- **Main identification idea.** Full-consideration discrete-choice models satisfy cross-price symmetry restrictions analogous to Slutsky symmetry. Limited consideration breaks these symmetries when product characteristics affect the probability that alternatives enter consideration sets.
- **Key contribution.** The paper shows that these asymmetric demand responses can constructively identify consideration probabilities and preferences in default-specific consideration, alternative-specific consideration, and hybrid models.
- **Validation.** The authors validate the approach in a lab experiment where the true consideration sets are known to the researcher but hidden from the estimator.
- **Application.** The paper applies the framework to Medicare Part D plan choice, where consumers are highly inertial and policymakers have proposed smart-default policies that reassign consumers to lower-cost plans while allowing opt-out.
- **Main empirical message.** Much of Part D inertia reflects limited consideration, not only utility-relevant switching costs. Conditional on paying attention, consumers are more responsive to plan attributes than a full-consideration conditional logit would suggest.
- **Welfare message.** Smart defaults can improve welfare if targeted to consumers with large potential monetary savings, but broad reassignment can have limited average benefits once acclimation and attention costs are accounted for.

2 Incremental contribution of the paper

- **Constructive identification without auxiliary consideration data.** Prior applied consideration-set models often used survey/search data or exclusion restrictions to learn what consumers considered. This paper shows that consideration probabilities can be recovered from choice data itself through cross-derivative asymmetries.
- **No need to exclude variables from either utility or attention.** Many previous models required some variables to affect attention but not utility, or utility but not attention. The paper allows the same observed attributes, including prices, to affect both utility and consideration.
- **Discrete-choice analog of Slutsky asymmetry.** The paper turns a theoretical insight—that inattention breaks symmetry—into an identification and estimation framework for applied demand.
- **General framework for DSC, ASC, and hybrid models.** The results cover both default-specific consideration models and alternative-specific consideration models, plus combinations of the two. This matters because real markets can feature both inertia/default choice and limited search among active consumers.
- **Preference-attention separation.** The paper explains how to avoid interpreting high choice probabilities as high utility when the product may simply be more likely to enter the consideration set.
- **Validation with known consideration sets.** The lab experiment is important because it provides a rare setting where the researcher can compare estimates against true consideration and preference parameters.
- **Policy use in smart defaults.** The framework shows that welfare evaluation of default policies depends critically on whether inertia reflects inattention or utility-relevant switching/acclimation costs.

- **Practical applied-econometrics contribution.** The authors show how the theory can be estimated by maximum likelihood or indirect inference and provide an empirical framework that can be applied to insurance, retirement plans, school choice, and online shopping.

Incremental contribution in one sentence. Abaluck and Adams-Prassl provide a way to identify and estimate limited consideration from the asymmetric demand responses already present in choice data, allowing researchers to separate utility from attention and to evaluate policies such as smart defaults without observing consumers' consideration sets.

3 Data and measurement

3.1 Lab experiment

- Participants made repeated choices from product-selection screens with Yale-branded merchandise.
- The researcher knew which subset of goods was shown or available in each round, but the estimation exercise deliberately ignores this information to mimic a standard field setting with unobserved consideration.
- The laboratory design has 10 possible goods and known structural parameters for attention/consideration.
- The purpose is validation: if the proposed method works, it should recover attention probabilities and preference parameters close to the known benchmark.

Interpretation: The lab experiment is not the main policy application. Its role is to show that asymmetric cross-price responses really can recover consideration and preference parameters when the truth is known.

3.2 Medicare Part D application

- The application uses Medicare Part D prescription drug plan choices, a setting with many plan options and strong inertia.
- The relevant choice set includes prescription drug insurance plans available to a beneficiary in a market-year.
- The paper focuses on returning beneficiaries, many of whom stay with their previous plan.
- The main policy counterfactual is a smart-default policy that switches consumers to lower-cost plans while allowing opt-out.

Interpretation: Part D is a natural setting because consumers face complex menus, high stakes, and defaults. Full-consideration models can interpret inertia as large switching costs; consideration-set models can instead attribute inertia to inattention.

3.3 Plan attributes

- Annual premium.
- Expected annual out-of-pocket costs.
- Variance of costs.
- Deductible.
- Donut-hole coverage.
- Average consumer cost sharing.

- Normalized plan quality rating.
- Number of top 100 drugs in the plan formulary.

Interpretation: These attributes enter both utility and consideration. This is essential: the paper does not rely on excluding premiums, quality, or other plan attributes from attention or utility.

3.4 Key attention objects

- **DSC component.** A consumer may be asleep and simply choose the default plan.
- **ASC component.** If awake, a consumer may actively consider only a subset of available alternatives.
- **Hybrid model.** Combines both channels: consumers may be asleep, or awake but still consider only a subset.

Interpretation: This hybrid structure lets the paper separate pure inertia from limited active search.

4 Models and empirical specifications

4.1 General consideration-set framework

Observed choice probabilities are a mixture over possible consideration sets:

$$s_j(p) = \sum_{C \in \mathcal{P}(j)} \pi_C(p) \tilde{s}_j(p | C),$$

where $s_j(p)$ is the observed probability of choosing product j , $\pi_C(p)$ is the probability that the consumer considers set C , and $\tilde{s}_j(p | C)$ is the probability that j is chosen conditional on C .

Interpretation: The observed market share is a mixture of two components: whether the alternative was considered and how attractive it is conditional on being considered. Standard discrete-choice models impose $\pi_C = 1$ for the full choice set and zero otherwise.

4.2 Daly-Zachary / full-consideration symmetry

Within a consideration set, latent choice probabilities satisfy symmetry of cross-price derivatives:

$$\frac{\partial \tilde{s}_j(p | C)}{\partial p_{j'}} = \frac{\partial \tilde{s}_{j'}(p | C)}{\partial p_j}.$$

Interpretation: In full-consideration random-utility models, cross-price responses are symmetric in the relevant Slutsky-like sense. If observed shares violate that symmetry, something outside full consideration is happening.

4.3 Identification from asymmetric responses

Limited consideration can make

$$\frac{\partial s_j(p)}{\partial p_{j'}} \neq \frac{\partial s_{j'}(p)}{\partial p_j}.$$

The paper shows that, in DSC/ASC/hybrid models, these asymmetries identify consideration probabilities.

Interpretation: The intuition is simple. If raising the price of a default wakes people up but lowering the price of a rival does not, the two cross-price effects are asymmetric. That asymmetry tells us about attention, not just preference.

4.4 Alternative-specific consideration model

In the lab validation, the probability that good j enters the consideration set is modeled as

$$\phi_{ijr} = Pr(\gamma_j + p_{irj}\gamma_p - \eta_{irj} > 0) = \frac{\exp(\gamma_j + p_{irj}\gamma_p)}{1 + \exp(\gamma_j + p_{irj}\gamma_p)}.$$

Interpretation: This ASC model lets price affect whether a product is considered. This is a major departure from models that impose consideration probabilities independent of price.

4.5 Hybrid Medicare Part D model

The hybrid model combines:

- a probability $\mu(\cdot)$ that the consumer is awake and does not automatically choose the default; and
- a probability $\phi(\cdot)$ that each nondefault alternative enters the consideration set conditional on being awake.

Interpretation: The model separates two questions: does the consumer wake up, and if awake, which plans does the consumer attend to?

4.6 Smart-default welfare simulations

The smart-default counterfactual reassigns consumers to lower-cost plans subject to thresholds and then lets them opt out according to the structural model. Welfare is computed under assumptions about attention costs and acclimation/switching costs.

Interpretation: Welfare depends on the mechanism behind inertia. If inertia is preference-based switching cost, forcing default changes can hurt. If inertia is attention-based, default changes can help.

5 Identification and empirical design

5.1 Core identification problem

The same observed choice can reflect utility or consideration. A plan with high share may be loved by consumers, or it may simply be salient/default. A conventional conditional logit cannot distinguish those mechanisms.

5.2 Cross-price asymmetry solution

The paper uses asymmetries in how demand responds to attribute changes. Full-consideration models impose symmetry. Consideration models can violate it because product attributes affect attention probabilities.

5.3 Why this is powerful

- It does not require survey data on consideration sets.
- It does not require excluding prices from either utility or attention.
- It applies to cross-sectional settings and does not require panel data for identification in the same way some alternatives do.
- It recovers consideration probabilities and preferences in widely used DSC/ASC/hybrid frameworks.

5.4 Validation and application logic

The lab experiment validates the method against known consideration sets. The Medicare application then uses the model where true consideration sets are unknown and where welfare counterfactuals are policy relevant.

5.5 Limitations

- The constructive identification results require structure on consideration-set formation.
- The model does not fully microfound why attention changes with product characteristics.
- Welfare depends on assumptions about acclimation costs and attention costs.
- Smart-default predictions are based on counterfactual simulation, not on an implemented policy experiment.

6 Section-by-section study guide

- **Introduction.** Explains why observed choices cannot distinguish utility from limited consideration and motivates Slutsky-asymmetry identification.
- **Related literature.** Positions the paper relative to consideration sets, rational inattention, and partial identification approaches.
- **Model and identification.** Develops the general framework, Daly-Zachary assumptions, DSC, ASC, and hybrid identification.
- **Lab validation.** Uses known consideration sets to verify that the method recovers attention and preferences better than conditional logit.
- **Medicare Part D.** Applies the hybrid model to plan choice and estimates attention, preference, and switching/acclimation parameters.
- **Smart defaults.** Simulates welfare effects under different reassignment thresholds and attention-cost assumptions.
- **Conclusion.** Emphasizes broad applicability to defaults, insurance, 401(k), school choice, and online shopping.

7 Tables and figures

7.1 Figure I: Lab Experiment Sample Product Selection Screen + Figure II: Product Fixed Effects in Attention

Read:

- Figure I shows the lab choice interface with a subset of Yale-branded products, prices, and a required delay before the next page.
- Figure II compares true product fixed effects in attention with estimates from the ASC model.
- The graph shows that the method can recover the broad pattern of which products are more likely to be considered.

Detailed interpretation: Figure I gives the reader an intuitive understanding of the lab experiment: consumers do not see all possible products at all times. Figure II then shows why the lab is useful: the true consideration parameters are known, so the authors can ask whether asymmetric demand responses recover them.

Economic message: The lab validates the paper’s main empirical claim before moving to Medicare, where consideration is unobserved.



FIGURE I

Lab Experiment: Sample Product Selection Screen

(a) Figure I: lab product selection screen.




FIGURE II

Product Fixed Effects in Attention: Truth versus ASC Model

(b) Figure II: product fixed effects in attention.

7.2 Table I: Experimental Data Estimation Results

Read:

- The conditional logit price coefficient is much too small in magnitude because it confounds high price with high consideration.
- The ASC model recovers more elastic price responses and separates utility from attention.
- The table compares conditional logit, ASC maximum likelihood, ASC indirect inference, and the true/conditional-on-consideration benchmark.

Detailed interpretation: Table I is the paper’s core validation table. It demonstrates that a full-consideration conditional logit can misread consideration as utility. If higher-priced products are more likely to be shown or considered, conditional logit incorrectly concludes consumers are less price-sensitive. The ASC model corrects this by estimating consideration directly.

Economic message: Ignoring consideration leads to biased demand elasticities and misleading welfare conclusions.

TABLE I
EXPERIMENTAL DATA ESTIMATION RESULTS

	Conditional logit (1)	ASC model		Conditional on consideration (4)
		MLE (2)	Indirect inf. (3)	
Utility				
Price (dollars)	−0.054*** (0.003)	−0.196*** (0.028)	−0.157** (0.029)	−0.173*** (0.004)
Product 1	−1.411*** (0.054)	1.465*** (0.539)	0.613 (0.406)	0.368*** (0.069)
Product 2	−1.955*** (0.069)	−0.065 (0.478)	0.133 (0.573)	−0.497*** (0.080)
Product 3	−1.627*** (0.059)	0.625 (0.476)	0.344 (0.403)	0.093 (0.073)
Product 4	−1.640*** (0.060)	0.629 (0.466)	0.443 (0.386)	0.088 (0.073)
Product 5	−1.447*** (0.056)	0.707 (0.478)	0.934*** (0.354)	0.306*** (0.070)
Product 6	−0.435*** (0.039)	−0.737*** (0.121)	−0.422 (0.252)	−0.581*** (0.045)
Product 7	−0.855*** (0.045)	−1.280*** (0.141)	−0.474 (0.234)	−1.075*** (0.051)
Product 8	−0.662*** (0.041)	−1.185*** (0.137)	−0.711*** (0.271)	−0.909*** (0.048)
Product 9	−0.316*** (0.038)	−0.561*** (0.118)	−0.706*** (0.184)	−0.405*** (0.044)
Attention				
Price (dollars)		0.137*** (0.017)	0.133*** (0.027)	0.15
Product 1		−2.872*** (0.177)	−2.636*** (0.359)	−2.5
Product 2		−2.674*** (0.288)	−2.914*** (0.544)	−2.5
Product 3		−2.695*** (0.209)	−2.599*** (0.360)	−2.5
Product 4		−2.704*** (0.205)	−2.703*** (0.389)	−2.5
Product 5		−2.592*** (0.204)	−3.042*** (0.389)	−2.5
Product 6		0.152 (0.192)	−0.464 (0.367)	0
Product 7		0.123 (0.292)	−0.925*** (0.345)	0
Product 8		0.258 (0.230)	0.476 (0.425)	0
Product 9		0.103 (0.176)	0.779 (0.532)	0

Table I: Experimental Data Estimation Results.

7.3 Table II: Sample Demographics and Plan Characteristics + Table III: Excess Sensitivity to Default Attributes in Switching Model

Read:

- Table II shows the Medicare Part D sample: average age is 75.5, 64.3 percent female, and 94.4 percent white.
- The average chosen plan has annual premiums of about \$415 and expected out-of-pocket costs of about

\$839.

- The average beneficiary faces about 11.9 options, and 93.7 percent remain inertial.
- Table III shows switching is excessively sensitive to attributes of the default plan, especially premiums, deductibles, and donut-hole coverage.

Detailed interpretation: These two tables motivate the hybrid model. Table II shows a complex environment with many plans and strong inertia. Table III shows that switching responds strongly to default-plan attributes, more than rival-plan attributes in ways a full-consideration model would not naturally predict.

Economic message: The Part D market is exactly where limited consideration matters: many alternatives, high inertia, and large potential savings.

	Mean	Std. dev.
Beneficiary characteristics:		
Age	75.5	8.6
Female	0.643	—
White	0.944	—
Plan characteristics of chosen plans		
Premiums	\$415	\$191
Out-of-pocket costs	\$839	\$713
Deductible	\$59	\$113
Donut hole coverage	0.121	—
Choice characteristics		
Options in Choice set	11.9	2.43
Inertial	93.7	—
No. beneficiary-plan-year	2,261,878	
No. beneficiary observations	79,286	

(a) Table II: sample demographics and plan characteristics.

	Share weighted	Lowest 3 plans	Lowest-cost plan
Annual premium (hundreds)	0.0801*** (0.0305)	0.0931*** (0.0348)	0.0914*** (0.0292)
Annual out-of-pocket costs (hundreds)	0.0057** (0.0028)	− 0.0051 (0.0048)	− 0.0096** (0.0046)
Variance of costs (millions)	− 0.0075** (0.0037)	0.0040 (0.0058)	− 0.0006 (0.0003)
Deductible (hundreds)	0.1203** (0.0545)	0.1242** (0.0501)	0.1219** (.0556)
Donut hole coverage	0.0827** (0.0410)	0.0926** (0.0376)	0.0909** (0.0396)
Average consumer cost sharing %	0.0186 (0.0141)	0.0234 (0.0161)	0.0210 (0.0143)
Normalized quality rating	0.0024 (0.0156)	0.0479 (0.1663)	0.0103 (0.0136)
# of top 100 drugs in formulary	0.1669 (0.0985)	0.0183 (0.0157)	0.1347 (0.1158)

Notes. This table reports coefficients from a panel data regression of an indicator for whether individual i switched at time t on attributes of the default plan, as well as the difference between default plan attributes and rival plan attributes for three different models of rival plan attributes. All models include individual and time fixed effects. In all cases, the reported coefficient is the coefficient on the default attribute conditioning on the difference between the default and rival plan attribute (α_d in equation (37)). The first column computes rival plan attributes as a share-weighted average of nondefault plans where the shares are computed for each (year, state, plan) using the choices of new beneficiaries (who are not included in this regression). The second column uses (unweighted) average attributes among the three lowest-cost plans. The third column uses attributes of the lowest-cost plan for rival plan attributes. The regression also includes an indicator for plans that are missing the # of top 100 drugs in formulary variable as well as an interaction of variance of costs and an indicator for individuals with no claims. Standard errors are in parentheses. *** denotes significance at the 1% level, ** at the 5% level, and * at the 10% level.

consider other available plans. A notable exception to this pattern

(b) Table III: excess sensitivity to default attributes.

7.4 Table IV: Conditional Logit and Hybrid Model

Read:

- Conditional logit implies a very large switching cost of about \$1,224.
- The hybrid model implies a much smaller acclimation/switching cost of about \$287.
- The model estimates separate utility, active-search probability $\mu(\cdot)$, and alternative-specific attention probability $\phi(\cdot)$.
- Mean attention probability is about 67.46 percent in the ASC component and about 15.09 percent for waking up in the DSC component.

Detailed interpretation: Table IV is the main structural table. The conditional logit forces inertia to be explained by utility-relevant switching costs. The hybrid model allows inertia to come from inattention, which sharply reduces the estimated acclimation cost and changes welfare implications.

Economic message: Separating attention from preferences turns a market that looks dominated by huge switching costs into one dominated by limited consideration.

TABLE IV
CONDITIONAL LOGIT AND HYBRID MODEL

	Conditional logit (1)	Hybrid model		
		Utility (2)	$\phi(\cdot)$ (3)	$\mu(\cdot)$ (4)
Annual premium (hundreds)	-0.4298*** (0.0103)	-0.6293*** (0.0331)	-1.118*** (0.0731)	0.0281** (0.0116)
Annual out-of-pocket costs (hundreds)	-0.1420*** (0.0131)	-0.0054 (0.0321)	-0.3797*** (0.0697)	0.0657*** (0.0206)
Variance of costs (millions)	0.1626*** (0.0183)	-0.1687*** (0.0492)	0.5501*** (0.0981)	1.5589*** (0.0301)
Deductible (hundreds)	-0.4108*** (0.021)	-0.6642*** (0.0552)	-0.0531 (0.1017)	0.2008*** (0.0483)
Donut hole coverage	0.5656*** (0.0609)	1.3579*** (0.2071)	10.1685*** (0.7432)	-0.0889 (0.1127)
Average consumer cost sharing %	-0.2209*** (0.0242)	-1.2484*** (0.0678)	1.3849*** (0.1474)	-0.3911*** (0.0663)
# of top 100 drugs in formulary	0.1243*** (0.0122)	0.0010 (0.0362)	0.4557*** (0.0483)	-0.1157*** (0.0371)
Normalized quality rating	0.0399*** (0.0021)	-0.0066 (0.0062)	0.1976*** (0.0181)	-0.0843*** (0.0043)
Default	5.2599*** (0.0224)	1.8052*** (0.1473)		
Constant			1.626*** (0.2126)	-1.853*** (0.0764)
Switching cost	\$1,224	\$287		
Mean attention probability			67.46%	15.09%

Notes. This table gives estimates from the conditional logit and “hybrid” version of the consideration set framework. Estimates in all models are the coefficients in the utility and attention equations (not marginal effects). The coefficients in the DSC component of the model, $\mu(\cdot)$, are the coefficients on the listed characteristics of the default good. The coefficients in the ASC component of the model, $\phi(\cdot)$, are the coefficients on the listed characteristics of good j on the likelihood of j being considered. The model also includes an indicator for plans that are missing the # of top 100 drugs in formulary variable as well as an interaction of variance of costs and an indicator for individuals with no claims. Standard errors are in parentheses. *** denote significance at the 1% level, ** at the 5% level, and * at the 10% level.

Table IV: Conditional Logit and Hybrid Model.

7.5 Figure III: Distribution of Savings + Table V: Welfare Impact of a Smart Default Policy

Read:

- Figure III shows the distribution of potential monetary savings from switching to the lowest-cost plan.
- Average savings are positive, and a nontrivial right tail can save more than \$1,000.
- Table V shows that defaulting everyone to the lowest-cost plan has small positive average welfare under the hybrid baseline, but results depend on attention-cost assumptions.
- Under full consideration, the policy has no effect because consumers opt out or switch back.

Detailed interpretation: Figure III shows why a smart default is tempting: many consumers leave money on the table. Table V shows why welfare analysis is subtle: if consumers face acclimation or attention costs, moving them automatically can create costs as well as savings.

Economic message: Smart defaults are not mechanically good. They help when attention failures dominate switching/acclimation costs.




FIGURE III

Distribution of Savings

is figure shows the distribution of potential monetary savings (annual

(a) Figure III: distribution of savings.

TABLE V
WELFARE IMPACT OF A SMART DEFAULT POLICY

	Market shares conditional on being switched			Attention cost			
	Smart default	Previous plan	Other plan	\$0	\$50	\$100	\$200
Conditional logit	1.000	0.000	0.000	\$0	\$0	\$0	\$0
Hybrid model	0.961	0.001	0.048	\$46	\$42	\$37	\$29
Direct effect on μ							
25%	0.946	0.002	0.053	\$39	\$32	\$25	\$12
50%	0.902	0.003	0.095	\$18	\$2	-\$14	-\$46
75%	0.853	0.005	0.142	-\$3	-\$30	-\$57	-\$111
100%	0.804	0.006	0.189	-\$25	-\$63	-\$101	-\$177
Proportion Switched by Policy: 89%							

Notes. This table shows overall welfare effects of a smart default policy where consumers are switched to the lowest-cost plan available in their market. Each row shows results with different assumptions about the direct effect of the smart default policy on the probability of saving attention, and each column shows

(b) Table V: welfare impact of a smart default policy.

7.6 Table VI: Welfare Effect of a Smart Default Policy: Restricted Reassignment to Those with Savings Greater than Switching Costs

Read:

- Table VI restricts reassignment to consumers whose monetary savings exceed a threshold ξ .
- When the policy targets those with larger savings, welfare gains are much larger.
- For the full sample, the hybrid-parameter welfare gain is about \$114, compared with much smaller gains under broad reassignment.
- Higher thresholds reduce the share switched but raise welfare among those switched.

Detailed interpretation: Table VI is the policy-design table. The main lesson is targeting: broad smart defaults may help only a little, but restricted reassignment toward consumers with large expected gains can generate much larger welfare improvements.

Economic message: The best default policy is not necessarily universal. Screening for large potential gains can preserve benefits while limiting unnecessary switching costs.

TABLE VI
WELFARE EFFECT OF A SMART DEFAULT POLICY: RESTRICTED REASSIGNMENT TO THOSE WITH SAVINGS GREATER THAN SWITCHING COSTS

	Market shares conditional on being switched			Attention cost			
	Smart default	Previous plan	Other plan	\$0	\$50	\$100	\$200
Mean welfare change: full sample							
Hybrid parameters				\$114	\$113	\$112	\$111
Direct effect on attention probability							
25%				\$109	\$107	\$106	\$102
50%				\$94	\$90	\$85	\$76
75%				\$79	\$71	\$64	\$49
100%				\$63	\$52	\$42	\$21
Proportion consumers switched by policy: 24%							
Mean welfare change: conditional on being switched							
Hybrid parameters	0.947	0.001	0.052	\$479	\$476	\$472	\$465
Direct effect on attention probability							
25%	0.923	0.002	0.076	\$458	\$451	\$444	\$429
50%	0.852	0.003	0.145	\$395	\$377	\$358	\$320
75%	0.778	0.005	0.217	\$330	\$299	\$267	\$205
100%	0.705	0.006	0.289	\$264	\$220	\$177	\$89

Notes. This table shows overall welfare effects of a smart default policy where consumers who will experience a monetary saving of at least ξ are switched to the lowest-cost plan available in their market. Each row shows results with different assumptions about the direct effect of the smart default policy on the probability of paying attention, while each column shows alternative assumptions about the cost of paying attention.

Table VI: Restricted reassignment to those with savings greater than switching costs.

8 Result map

- **Theory result 1:** full consideration implies cross-price symmetry; limited consideration can violate it.
- **Theory result 2:** in DSC/ASC/hybrid models, these asymmetries identify consideration probabilities.
- **Validation result:** the ASC model recovers the lab attention process better than conditional logit.
- **Medicare result:** Part D consumers are highly inertial and often attend to only a few plans.
- **Structural result:** the hybrid model greatly reduces estimated switching/acclimation costs relative to conditional logit.
- **Policy result:** broad smart defaults have small benefits, but targeted smart defaults can create substantial welfare gains.

9 Short summary

- The paper identifies limited consideration using asymmetric demand responses.
- It does not require auxiliary consideration data or exclusion restrictions.
- Cross-price asymmetries play the role of a discrete-choice Slutsky test.
- The lab experiment validates the method when consideration sets are known.
- Conditional logit confounds consideration with utility and understates price sensitivity.
- The Medicare Part D application uses a hybrid model with sleeping/default behavior and active limited consideration.
- Most inertia in Part D reflects inattention, but there are also nontrivial utility-relevant acclimation costs.
- Smart defaults should be targeted toward beneficiaries with large potential savings.

10 Conclusion

10.1 Core conclusion

Abaluck and Adams-Prassl show that consideration sets can be identified from choice data by exploiting violations of cross-price symmetry. This turns a major obstacle in limited-consideration demand estimation into a source of identification.

10.2 Economic intuition

If all consumers consider all goods, demand responses should be symmetric in a Slutsky-like sense. If consumers ignore some goods, a price change can affect utility and attention at the same time. The attention component breaks symmetry, and the pattern of asymmetries reveals what consumers consider.

10.3 Incremental contribution revisited

The paper’s incremental contribution is methodological and applied. Methodologically, it gives constructive identification of consideration probabilities without extra consideration data. Applied, it shows that policy evaluation can change dramatically when inertia is interpreted as limited attention rather than only switching costs.

10.4 Policy implications

For Medicare Part D, smart defaults are potentially valuable because many consumers do not actively evaluate all plans. But the policy should be targeted. Reassigning everyone can impose unnecessary acclimation and attention costs, while reassigning consumers with large potential savings produces larger welfare gains.

10.5 Limitations

The framework uses structure from DSC, ASC, or hybrid consideration models. It does not provide a full psychological or rational-inattention microfoundation for why attention changes. Counterfactuals require assumptions about how attention responds under new policies.

10.6 Takeaway

Consumer choice data contain more information than standard discrete-choice models use. Asymmetries in demand responses can reveal the hidden consideration process, allowing researchers to separate what consumers like from what consumers notice.