Notes: Ifrach and Weintraub (REStud, 2017)

A Framework for Dynamic Oligopoly in Concentrated Industries

Contents

1 Big picture **3**

2 Data and measurement (key objects) **3**

2.1 Core state variables and objects 4
2.2 What firms observe versus what they track 4
2.3 Entry, exit, and dynamic actions 4
2.4 Key empirical and calibration objects in the numerical sections 5
2.5 Key output objects the paper studies 5

3 Models and empirical specifications (all equations + interpretation) **5**

3.1 Baseline dynamic oligopoly model 5
3.2 Markov perfect equilibrium (MPE) 6
3.3 Dominant/fringe decomposition 6
3.4 Moment-based strategies 6
3.5 Why the reduced state is not really Markov 7
3.6 Perceived transition kernel 7
3.7 Perceived single-period profit function 8
3.8 Definition of moment-based Markov equilibrium (MME) 8
3.9 Algorithm to compute MME 8
3.10 Relation to other equilibrium concepts 9
3.11 Two benchmark profit-function examples 9
3.12 Numerical comparison of MME and MPE 10
3.13 Large-scale application: the beer industry 10
3.14 MME as approximation to MPE and as a behavioral model 10
3.15 Robust upper bound for the deviation value 11

4 Identification and instruments (what makes the estimates credible) **11**

4.1 What disciplines the reduced state 11
4.2 Why the approximation is plausible 12
4.3 What makes the quantitative exercise credible 12
4.4 Why the beer application is informative 12
4.5 Main limitations 12
4.6 Relation to the literature 13

5 Tables and figures **13**
5.1 Table 1: default parameters for the benchmark numerical models 13
5.2 Table 2: how closely does MME approximate MPE? 14
5.3 Figure 1 and Table 3: the beer industry application 14
5.4 Figure 2: error bounds for unilateral deviations 15

6 Short summary **15**

7 Conclusion **16**
7.1 Core conclusion (what the paper establishes) 16
7.2 Economic intuition 16
7.3 Takeaway 16

1 Big picture

- **Question.** How can one study dynamic oligopoly in industries with a few dominant firms and many fringe firms when exact computation of Markov perfect equilibrium (MPE) is infeasible?
- **Setting.** The paper works in Ericson–Pakes style dynamic oligopoly models with heterogeneous firms, entry, exit, investment, and possibly aggregate shocks. The difficulty is that the full industry state becomes enormous once there are many fringe firms.
- **Core computational problem.** In standard MPE, a firm's strategy depends on its own state and the full distribution of rivals. With many fringe firms, the state space explodes and dynamic programming becomes computationally impractical.
- **Main idea.** Replace the full fringe-state distribution with a low-dimensional summary. Firms keep track of: (i) their own state, (ii) the detailed state of dominant firms, and (iii) a small number of moments of the fringe distribution.
- **New equilibrium concept.** The authors define *moment-based Markov equilibrium* (MME): an equilibrium in these reduced-state strategies, together with a perceived Markov transition kernel over the reduced state.
- **Why a perceived kernel is needed.** Moments are generally *not* sufficient statistics for future industry dynamics. Different fringe distributions can generate the same current moment but imply different future paths. So the reduced state is not truly Markov, and the paper approximates it with a perceived Markov law of motion.
- **Main methodological contribution.** The paper gives a practical way to compute dynamic equilibria in concentrated industries with hundreds of firms, while still allowing fringe firms to discipline dominant firms and allowing firms to move between fringe and dominant tiers.
- **Main validation exercise.** In industries small enough that exact MPE can still be computed, MME closely approximates MPE for two benchmark models: a quality-ladder model and a capacity-competition model.
- **Main applied illustration.** In a calibrated beer-industry application, increasing returns to advertising generate a concentrated market structure with a few dominant firms and many small fringe firms. Explicitly modelling the fringe matters because fringe firms discipline dominant firms' investment incentives.
- **Main theoretical support for approximation.** The paper defines the *value of the full information deviation* and derives a computationally tractable upper bound for it using robust dynamic programming ideas.
- **Economic takeaway.** The paper is really about how to keep the strategic richness of dominant-versus-fringe industries while compressing the state space enough to make dynamic analysis possible.

2 Data and measurement (key objects)

This is mainly a methodological paper, so the core "data" are not reduced-form estimating moments in the usual empirical sense. The central objects are model primitives, low-dimensional state summaries, and then calibration targets in the numerical examples and beer application.

2.1 Core state variables and objects

- **Firm state.** Each incumbent firm i has an individual state $x_{it} \in X$, where X is finite. Depending on the application, this state can mean product quality, productivity, capacity, goodwill, or some other payoff-relevant stock.
- **Industry state.** The full industry state is the combination of:
  - the fringe-firm distribution f_t over states,
  - the detailed vector of dominant firms' states d_t,
  - the aggregate shock z_t.
- **Dominant versus fringe firms.** The market structure is explicitly split into a small dominant tier and a large fringe tier. A firm becomes dominant when its state crosses a threshold or satisfies another application-specific rule.
- **Moments of the fringe state.** A low-dimensional function

  $$\theta_t = \theta(f_t)$$

  summarizes the fringe distribution. Examples include un-normalized moments, quantiles, total fringe capacity, or the number of fringe firms close to becoming dominant.

2.2 What firms observe versus what they track

- In the underlying complete-information game, the full industry state is in principle observable.
- In the MME approximation, firms are assumed to track only:
  - their own state x_{it},
  - dominant-firm states d_t,
  - aggregate shock z_t,
  - fringe moments θ_t.
- This is meant to reflect the idea that dominant firms are visible and strategically important, whereas tracking every small fringe rival is unrealistic and computationally prohibitive.

2.3 Entry, exit, and dynamic actions

- **Exit.** Incumbents observe private sell-off values ϕ_{it} and exit if the outside option exceeds continuation value.
- **Entry.** Potential entrants observe private entry costs κ_{it} and enter if expected value exceeds the cost.
- **Investment.** Firms choose dynamic actions ι_{it} to improve their states over time.
- **Aggregate shock.** A finite-state Markov process z_t shifts profitability.

2.4 Key empirical and calibration objects in the numerical sections

- In the small-scale comparison section, the authors calibrate two stylized models using plausible parameter values rather than matching a specific data set one-to-one.
- In the beer application, the paper calibrates a dynamic advertising model using evidence from the U.S. beer industry and related advertising studies.
- The large-scale beer experiments use:
  - $N = 200$ firms,
  - 29 individual states,
  - at most $D = 3$ dominant firms,
  - alternative returns-to-advertising cases: DRA–DRA, DRA–CRA, and CRA–IRA.

2.5 Key output objects the paper studies

- long-run total investment,
- producer surplus,
- consumer surplus,
- social welfare,
- concentration ratios,
- the long-run firm-size distribution,
- and the bound on the value of deviating from MME to a full-state Markov best response.

3 Models and empirical specifications (all equations + interpretation)

3.1 Baseline dynamic oligopoly model

Time is infinite and discrete. Firms have individual states, can enter and exit, and invest to improve their states. The full industry state is

$$s_t = (\bar{s}_t, z_t),$$

where $\bar{s}_t$ is the distribution of firms across individual states and z_t is the aggregate shock.

A surviving incumbent that invests ι transitions according to

$$Q[x' \mid x, \iota] = \Pr(x_{i,t+1} = x' \mid x_{it} = x, \iota_{it} = \iota),$$

and pays investment cost $c(\iota, x)$.

The value of a firm that deviates to strategy μ' while competitors use μ and entrants use λ is

$$V(x, s \mid \mu', \mu, \lambda) = \mathbb{E}^{\mu',\mu,\lambda}\left[\sum_{k=t}^{\tau_i} \beta^{k-t}\big(\pi(x_{ik}, s_k) - c(\iota_{ik}, x_{ik})\big) + \beta^{\tau_i - t}\phi_{i,\tau_i} \;\Big|\; x_{it} = x, s_t = s\right].$$

Interpretation. This is the standard Ericson–Pakes structure: dynamic profits depend on current operating profits, investment costs, and the option to exit.

3.2 Markov perfect equilibrium (MPE)

A standard MPE is a strategy pair (μ, λ) such that incumbent firms best respond given the full industry state and entrants use a cutoff rule satisfying

$$\sup_{\mu' \in M} V(x, s \mid \mu', \mu, \lambda) = V(x, s \mid \mu, \lambda),$$

and

$$\lambda(s) = \beta \, \mathbb{E}^{\mu,\lambda}\big[V(x_e, s_{t+1} \mid \mu, \lambda) \mid s_t = s\big].$$

Interpretation. MPE is the exact equilibrium benchmark. The problem is that the state space grows explosively with the number of fringe firms and possible individual states.

3.3 Dominant/fringe decomposition

The paper exploits a market structure with a few large firms and many small ones.

- D_t is the set of dominant firms and F_t is the set of fringe firms.
- The fringe state is a distribution f_t over X_f.
- The dominant state is a low-dimensional ordered vector $d_t \in S_d$.

The complete state is then re-written as

$$s_t = (f_t, d_t, z_t).$$

Interpretation. This decomposition is economically motivated and computationally useful. Firms still react exactly to dominant firms, while the fringe is compressed into a distribution.

3.4 Moment-based strategies

Define a low-dimensional summary of the fringe state:

$$\theta_t = \theta(f_t).$$

The *moment-based state* is

$$\hat{s}_t = (\theta_t, d_t, z_t),$$

with state space $\hat{S} = S_\theta \times S_d \times Z$.

A moment-based investment strategy is

$$\iota_{it} = \iota(x_{it}, \hat{s}_t),$$

and a moment-based exit strategy is a cutoff rule

$$\text{exit if } \phi_{it} \geq \rho(x_{it}, \hat{s}_t).$$

Entrants use a cutoff rule

$$\text{enter if } \kappa_{it} \leq \lambda(\hat{s}_t).$$

Interpretation. MME compresses the informational burden of the fringe tier. The paper stresses that firms always keep at least those moments needed to compute or approximate the single-period profit function.

3.5 Why the reduced state is not really Markov

The problem is that the same moment can be generated by many different fringe distributions. For example, a first moment equal to 10 could come from one fringe firm of size 10 or ten fringe firms of size 1. These two full states can have very different next-period laws of motion.

Interpretation. This is the key conceptual difficulty. Once we aggregate the fringe state into moments, we generally lose the Markov property.

3.6 Perceived transition kernel

To repair the loss of the Markov property, the paper defines a *perceived* Markov transition kernel over the reduced state.

One nonparametric version uses observed transitions in long simulations:

$$\hat{P}^{\mu,\lambda}[\theta' \mid \hat{s}] = \limsup_{T\to\infty} \frac{\sum_{t=1}^{T} \mathbf{1}\{\hat{s}_t = \hat{s},\ \theta_{t+1} = \theta'\}}{\sum_{t=1}^{T} \mathbf{1}\{\hat{s}_t = \hat{s}\}}.$$

A more restrictive but faster parametric version assumes a law of motion such as

$$\theta_{t+1} = G(\theta_t; \xi(d_t, z_t)),$$

for example a linear rule

$$\theta_{t+1} = \xi_0(d_t, z_t) + \xi_1(d_t, z_t)\theta_t.$$

The joint perceived kernel combines:

- the transition of the firm's own state,
- the transition of current dominant firms,
- the aggregate shock,
- the perceived transition of moments,
- and perceived fringe-to-dominant transitions.

Interpretation. The paper does *not* claim the reduced state is truly Markov. It claims that one can construct a useful Markov approximation to it.

3.7 Perceived single-period profit function

Because firms using moment-based strategies do not track the full fringe state, the paper also defines a perceived single-period profit function over the reduced state. One specification is the long-run average primitive profit conditional on the reduced state:

$$\hat{\pi}_{\mu,\lambda}(x, \hat{s}) = \limsup_{T\to\infty} \frac{\sum_{t=1}^{T} \pi(x, s_t)\mathbf{1}\{\hat{s}_t = \hat{s}\}}{\sum_{t=1}^{T} \mathbf{1}\{\hat{s}_t = \hat{s}\}}.$$

The paper then imposes a practically important assumption: the single-period expected profit depends on the fringe only through a low-dimensional vector of moments,

$$\pi(x_{it}, \tilde{\theta}_t, d_t, z_t).$$

Interpretation. This is what makes the state compression economically meaningful. If current profits themselves can be approximated with a few fringe moments, then using those same moments in strategies is natural.

3.8 Definition of moment-based Markov equilibrium (MME)

Given the perceived kernel, define the perceived value function

$$\hat{V}(x, \hat{s} \mid \mu', \mu, \lambda) = \mathbb{E}^{\mu',\mu,\lambda}\left[\sum_{k=t}^{\tau_i} \beta^{k-t}\big(\hat{\pi}(x_{ik}, \hat{s}_k) - c(\iota_{ik}, x_{ik})\big) + \beta^{\tau_i - t}\phi_{i,\tau_i} \;\middle|\; x_{it} = x, \hat{s}_t = \hat{s}\right].$$

An MME is a pair (μ, λ) such that:

$$\sup_{\mu' \in \hat{M}} \hat{V}(x, \hat{s} \mid \mu', \mu, \lambda) = \hat{V}(x, \hat{s} \mid \mu, \lambda) \qquad \forall (x, \hat{s}) \in X \times \hat{S},$$

$$\lambda(\hat{s}) = \beta\, \mathbb{E}^{\mu,\lambda}\big[\hat{V}(x_e, \hat{s}_{t+1} \mid \mu, \lambda) \mid \hat{s}_t = \hat{s}\big] \qquad \forall \hat{s} \in \hat{S}.$$

Interpretation.

- This is the exact analogue of MPE, but in the reduced state space.
- If $\theta(f) = f$ and the perceived kernel equals the true kernel, MME collapses to MPE.
- So MME should be read as a reduced-state equilibrium concept, not a completely different game.

3.9 Algorithm to compute MME

The paper proposes a best-response iteration algorithm.

Algorithm 1: equilibrium solver

1. Start from a candidate strategy profile (μ, λ) and an initial industry state.
2. Simulate a long sample path under the candidate strategies.
3. Use that sample path to estimate the frequency of reduced states and the perceived transition kernel.

4. Solve the single-agent dynamic program to obtain incumbent best-response moment-based strategies.
5. Update the entrant cutoff rule using the implied continuation values.
6. Smooth the update and iterate until the distance between current strategies and best responses is small.

Interpretation.

- The algorithm is simulation-based.
- The expensive object is not the entire full-state transition law but the perceived reduced-state law.
- The paper also describes a faster real-time stochastic version akin to Pakes–McGuire style updating.

3.10 Relation to other equilibrium concepts

- MME is related to restricted experienced-based equilibrium (REBE) in Fershtman and Pakes, because both limit the information firms actively use.
- It generalizes partially oblivious equilibrium (POE) of Benkard, Jeziorski, and Weintraub by allowing firms to keep track not only of dominant firms but also of fringe moments, and by letting dominant status arise endogenously.
- It is in the spirit of Krusell–Smith style aggregation, but applied to dynamic oligopoly rather than macro heterogeneous-agent models.

3.11 Two benchmark profit-function examples

3.11.1 Example 6.1: quality ladder

Consumers choose among differentiated products with logit demand, where higher-quality firms are more attractive. Under a simplification in which fringe firms price as if they have no market power, profits depend on the fringe only through one moment,

$$\tilde{\theta}_t = \sum_{y \in X_f} y^{\alpha_1} f_t(y).$$

Interpretation. This moment summarizes the total competitive pressure exerted by the fringe on market shares.

3.11.2 Example 6.2: capacity competition

Firms compete in quantities with capacity constraints. Under the approximation that fringe firms are small and produce at full capacity, profits depend on the fringe only through total fringe capacity,

$$\tilde{\theta}_t = \sum_{x \in X_f} f_t(x)\, \bar{q}(x).$$

Interpretation. Total installed capacity in the fringe is a natural aggregate object for firms to track.

3.12 Numerical comparison of MME and MPE

The paper compares MME and exact MPE in small state spaces with $|X| = 9$ and $N = 5, 6, 7$ firms.

- In the quality-ladder model, most long-run indicators are approximated within roughly 3%–5%.
- In the capacity model, approximation errors are also generally small, though consumer surplus can deviate by around 11%–12% in some cases.
- MME computation takes minutes, whereas MPE can take many hours.

The paper emphasizes four sources of approximation error:

- fringe firms in MME do not strategically deter rivals until they are monitored as dominant,
- profits are approximated with a small set of moments,
- moments may not be sufficient statistics,
- the number of dominant slots may be too small.

Important implementation lesson. Adding more dominant slots often improves the approximation much more than adding one extra fringe moment. That is because dominant-slot scarcity can distort incentives of firms close to becoming dominant.

3.13 Large-scale application: the beer industry

The beer application calibrates a dynamic advertising model in which a firm's state is goodwill and market share satisfies

$$\sigma(x_{it}, s_t) \propto (x_{it})^{\alpha_1}(Y - p_{it})^{\alpha_2}.$$

Firms invest in advertising to raise goodwill. They compete in prices every period. The key parameter is the return to advertising, α_1.

The paper considers three cases:

- DRA–DRA: fringe and dominant firms both have decreasing returns to advertising,
- DRA–CRA: fringe firms have decreasing returns and dominant firms constant returns,
- CRA–IRA: fringe firms have constant returns and dominant firms increasing returns.

Interpretation. Increasing returns to advertising for dominant firms create an endogenous force toward concentration and long-lived dominance.

3.14 MME as approximation to MPE and as a behavioral model

The paper defines the *value of the full information deviation*

$$\Delta_{\mu,\lambda}(x, s) = \sup_{\mu' \in M} V(x, s \mid \mu', \mu, \lambda) - V(x, s \mid \mu, \lambda).$$

Interpretation. This is the gain from unilaterally deviating from the MME strategy to the best full-state Markov strategy, taking competitors' MME behavior as given.

If $\Delta_{\mu,\lambda}(x,s)$ is small, then MME is appealing as a reduced-information behavioral rule, not only as a computational approximation.

3.15 Robust upper bound for the deviation value

To bound the deviation value, define the consistency set for a reduced state $\hat{s} = (\theta, d, z)$:

$$S_f(\hat{s}) = \{f \in S_f : \theta(f) = \theta\}.$$

The robust Bellman operator lets the deviator choose the best full fringe state consistent with the observed moment:

$$\begin{aligned}(T_R\hat{V})(x,\hat{s}) = \sup_{\iota \in I,\ \rho \geq 0} \sup_{f \in S_f(\hat{s})} \Big\{ &\pi(x,\hat{s}) + \mathbb{E}\big[\phi \mathbf{1}\{\phi \geq \rho\} \\ &+ \mathbf{1}\{\phi < \rho\}\big(-c(x,\iota) + \beta\mathbb{E}[\hat{V}(x_{i,t+1},\hat{s}_{t+1}) \mid x_t = x, s_t = (f,d,z),\iota]\big)\big]\Big\}.\end{aligned}$$

Theorem 8.1 states that if $\hat{V}^*$ is the fixed point of this operator, then for every full state consistent with $\hat{s}$,

$$V^*(x,s) \leq \hat{V}^*(x,\hat{s}).$$

Hence,

$$\hat{\Delta}(x,s) = \hat{V}^*(x,\hat{s}) - V(x,s)$$

provides an upper bound on the true value of deviating to a full-state best response.

Interpretation.

- The robust bound resolves missing fringe information in the most favorable way for the deviator.
- So it is conservative: if even that bound is small, then MME is hard to improve upon.
- The paper further shows that computation becomes feasible when the fringe is large, because one-step fringe transitions are close to deterministic by a law of large numbers.

4 Identification and instruments (what makes the estimates credible)

This paper is not an IV paper in the standard empirical sense. "Identification" here means: what disciplines the reduced-state approximation, what makes the quantitative exercise credible, and what features of the model are actually being pinned down.

4.1 What disciplines the reduced state

- The chosen moments are not arbitrary summary statistics. The paper starts from moments that directly determine or closely approximate the single-period profit function.

- In the quality-ladder model, the natural moment summarizes fringe competitive pressure in demand.
- In the capacity model, total fringe capacity is the natural sufficient approximation object for profits.

Why this matters. If the moment already captures the main way the fringe matters for current profits, it is a good first candidate for also approximating continuation values.

4.2 Why the approximation is plausible

- In concentrated industries, firms plausibly care much more about dominant firms than about the exact identity of every small fringe rival.
- A single small fringe firm's effect on future aggregate fringe moments is also small when the fringe is large.
- As the number of fringe firms grows, one-step transitions of the aggregate fringe state become closer to deterministic by averaging out idiosyncratic shocks.

4.3 What makes the quantitative exercise credible

- The paper validates MME against exact MPE in environments small enough to solve exactly.
- It uses multiple economic indicators—investment, producer surplus, consumer surplus, and total welfare—rather than only strategy distances.
- It investigates where approximation errors come from and how they respond to practical design choices such as the number of dominant slots, added moments, and beliefs outside the recurrent class.
- It then moves to a setting where exact MPE is impossible, which is exactly where the new method is useful.

4.4 Why the beer application is informative

- The beer industry is a classic concentrated-industry case: many small firms, a few dominant leaders, and a widely discussed role for advertising as an endogenous sunk cost.
- The application is not pitched as a full structural estimate of the beer industry. Its role is to demonstrate how the method can endogenize market structure in a large dynamic model.

4.5 Main limitations

- Moments generally are not sufficient statistics, so MME remains an approximation.
- Perceived transitions outside the recurrent class are arbitrary and can matter, especially in smaller industries.
- The robust deviation bound can be loose, not just conservative.

- The paper focuses on "capital accumulation" games without investment spillovers; broader strategic environments would require further work.
- The beer application is illustrative rather than a complete empirical fit.

4.6 Relation to the literature

- Relative to exact EP computation, the paper sacrifices state richness in the fringe tier to gain tractability.
- Relative to oblivious equilibrium, it is built for concentrated industries and retains detailed dominant-firm states.
- Relative to partially oblivious equilibrium, it adds fringe moments and endogenizes dominant status.
- Relative to Krusell–Smith style aggregation, it applies moment compression to dynamic oligopoly and develops an equilibrium concept plus error bounds tailored to this setting.

5 Tables and figures

5.1 Table 1: default parameters for the benchmark numerical models

Read: This table sets the baseline parameter values for the quality-ladder and capacity-competition experiments.

- For the quality-ladder model, benchmark values include $m = 60$, marginal cost $d = 0.5$, demand parameters $\alpha_1 = 1$, $\alpha_2 = 0.5$, income $Y = 1$, and investment cost $c = 0.65$.
- For the capacity model, benchmark values include $m = 2$, $q_i = 0$, $q_s = 2$, $e = 50$, $f = 10$, and $c = 4$.
- Shared parameters include discount factor $\beta = 0.925$, depreciation probability $\delta = 0.7$, success parameter $a = 1$, and nine individual states.

Why it matters. These are not meant to identify one particular industry. They are benchmark economies used to compare MME against exact MPE.

TABLE 1

Default parameters for quality ladder model and capacity model experiments

Parameters quality ladder model	m	d	α_1	α_2	Y	c
Value	60	0.5	1	0.5	1	0.65
Parameters capacity model	m	$\bar{q}_i$	$\bar{q}_s$	e	f	c
Value	2	0	2	50	10	4
Parameters shared both models	β		δ	a	$\mathcal{X}$	
Value	0.925		0.7	1	$\{1,\ldots,9\}$	

TABLE 2

Comparison of MPE and MME indicators

Instance		Number of firms	Long-run statistics: Total Inv.	Prod. Surp.	Cons. Surp.	Soc. Welf.
Quality ladder model	High inv. cost	$N=5$	7.9	5.0	1.7	2.1
		$N=6$	3.1	4.4	2.5	2.7
		$N=7$	1.7	4.0	3.3	3.4
	Low inv. cost	$N=5$	7.8	5.0	1.7	2.1
		$N=6$	3.4	4.2	2.4	2.6
		$N=7$	1.4	3.7	3.1	3.1
Capacity competition model	High inv. cost	$N=5$	0.1	2.3	5.1	2.8
		$N=6$	2.8	2.8	7.8	3.7
		$N=7$	5.0	3.5	11.6	5.1
	Low inv. cost	$N=5$	4.6	1.5	5.7	2.7
		$N=6$	1.7	1.5	7.6	3.3
		$N=7$	1.0	1.7	11.9	4.7

Long-run statistics computed simulating industry evolution over 10^6 periods. Numbers presented are percentage differences of the economic indicators under MPE and MME.

5.2 Table 2: how closely does MME approximate MPE?

Read: This table reports percentage differences between MME and MPE in long-run total investment, producer surplus, consumer surplus, and social welfare.

- In the quality-ladder model, most errors are in the 1%–5% range.
- In the capacity model, most errors are still moderate, but consumer surplus can deviate more, reaching roughly 11%–12% in the harder cases.
- Computation time differs dramatically: MME can be solved in minutes, while MPE can take many hours even in these small state spaces.

Interpretation. The table is the paper's central credibility check. It shows that the reduced-state method reproduces the economically relevant objects of the exact equilibrium reasonably well.

5.3 Figure 1 and Table 3: the beer industry application

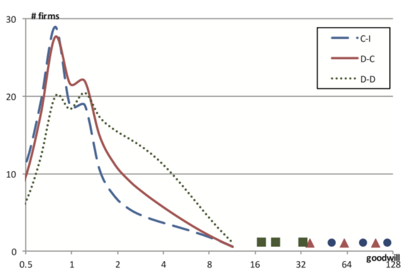

FIGURE 1

Long-run average size distribution (un-normalized) of firms in log-scale. The curve lines represent fringe firms for the different levels of returns to advertising. The squares represent dominant firms for D-D, the triangles for D-C, and the circles for C-I

TABLE 3
Average industry statistics

Industry averages	CRA-IRA	DRA-CRA	DRA-DRA
Concentration ratio C_2	0.42	0.36	0.13
Concentration ratio C_3	0.53	0.44	0.17
First moment (normalized by # fringe)	1.09	1.22	1.74
Active fringe firms (#)	147	166	180
Active dominant firms (#)	3	3	2.3
Size incumbent fringe (goodwill)	1.1	1.4	2
Size dominant (goodwill)	83.4	64.9	24.1
Entrants per period (#)	16.5	12.3	7.6
Time in industry fringe	12.9	16	26.2
Time as dominant	1536	360	21

Read Figure 1. The figure plots the long-run distribution of goodwill across firms under three returns-to-advertising specifications.

- DRA–DRA yields a relatively diffuse industry with weakly separated dominant firms.
- DRA–CRA and CRA–IRA produce much more skewed size distributions, with clear dominant firms and a smaller fringe.

Read Table 3. Key numbers are:

- concentration ratio C_2: 0.13 in DRA–DRA, 0.36 in DRA–CRA, and 0.42 in CRA–IRA,
- concentration ratio C_3: 0.17, 0.44, and 0.53 respectively,
- average number of active fringe firms: 180, 166, and 147,
- dominant-firm size: 24.1, 64.9, and 83.4,
- fringe time in industry: 26.2, 16, and 12.9,
- time as dominant: 21, 360, and 1536.

Interpretation.

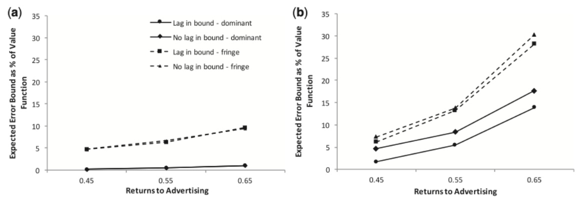

FIGURE 2

Error bound (a) No entry/exit; (b) entry/exit

- Higher returns to advertising make market structure more concentrated and dominance more persistent.
- The fringe becomes smaller, shorter-lived, and more strongly disciplined.
- This is exactly the kind of large-scale dynamic market-structure question the method is built to study.

Extra lesson from this section. If one ignores fringe firms and solves an oligopoly among dominant firms only, dominant firms invest much less. So explicitly modelling the fringe matters for counterfactual analysis.

5.4 Figure 2: error bounds for unilateral deviations

Read: This figure reports the robust upper bound on the gain from deviating from MME to a full-information Markov strategy in simplified beer-like models.

- Bounds are smaller when the fringe is more homogeneous.
- Bounds are also lower in the no-entry/no-exit model than in the entry/exit model because linear transition approximations work better in the simpler environment.
- Adding a lagged fringe state tightens the bound.

Interpretation. The bound is sometimes tight and sometimes loose, but it is still a practically useful diagnostic: if the bound is small, the MME approximation is hard to beat; if it is large, the user should enrich the reduced state.

6 Short summary

- The paper proposes **moment-based Markov equilibrium (MME)** for dynamic oligopoly in concentrated industries with a few dominant firms and many fringe firms.
- Firms track their own state, dominant-firm states, and a few moments of the fringe distribution instead of the full fringe distribution.
- Because moments are generally not sufficient statistics, the paper introduces a **perceived Markov transition kernel** over the reduced state.
- The method is validated against exact MPE in small models and works well on key long-run outcomes.

- In a beer-industry application, stronger returns to advertising generate much more concentrated industries with long-lived dominant firms.
- The paper also derives a **robust upper bound** on the gain from deviating to a full-information strategy, giving users a way to judge whether the state compression is good enough.

7 Conclusion

7.1 Core conclusion (what the paper establishes)

The paper establishes that one can study rich dynamic oligopoly problems in concentrated industries without solving the full Markov perfect equilibrium. The key is to preserve exact detail where it matters most—dominant firms—and compress the fringe into a few strategically relevant moments. This yields a tractable equilibrium concept, a workable computational algorithm, and a practical error diagnostic.

7.2 Economic intuition

The economic intuition is simple but powerful:

- dominant firms are few, visible, and strategically central,
- fringe firms are many, individually small, and mostly matter through aggregates,
- but the fringe still matters collectively, so it cannot simply be ignored.

MME keeps precisely this structure. It therefore sits between two extremes: exact MPE, which is often impossible to compute, and overly coarse approximations that eliminate the strategic role of the fringe.

7.3 Takeaway

If the research question is about dynamic market structure, policy counterfactuals, or dominant-versus-fringe interaction in industries with many firms, this paper's main lesson is: *use low-dimensional state aggregation, but validate it carefully.* MME is valuable not because it is exact, but because it is structured, interpretable, computationally feasible, and testable.