

Notes: Jones (RES, 2009)

The Burden of Knowledge and the “Death of the Renaissance Man”: Is Innovation Getting Harder?

Contents

1	Big picture	2
2	Incremental contribution of the paper	2
3	Data and measurement	3
3.1	Patent data	3
3.2	Inventor histories and age data	3
3.3	Main outcome variables	4
3.4	Why solo inventors matter for specialization	4
4	Empirical specifications and evidence	4
4.1	Time-series regressions	4
4.2	Cross-sectional regressions using knowledge depth	5
5	The model	6
5.1	Knowledge space	6
5.2	Idea arrival and crowding	6
5.3	Team formation	6
5.4	Equilibrium implications	7
5.5	Growth equation	7
6	Identification and empirical design	7
6.1	What the empirical design can identify	7
6.2	Key empirical challenges	8
6.3	Why the evidence is still useful	8
7	Tables and figures	8
7.1	Table 1 and Figure 1: motivating age trends and basic patent trends	8
7.2	Figure 2 and Table 2: data window and teamwork trend	9

7.3	Tables 3 and 4: age at first innovation and specialization over time	10
7.4	Tables 5 and 6: cross-field patterns and tree-size evidence for teamwork	10
7.5	Tables 7 and 8: tree size, age, and specialization	11
7.6	Figure 3 and Table B.1: model geometry and age-data selection	12
7.7	Tables B.2 and B.3: appendix checks on the age sample	13
8	Short summary	14
9	Conclusion	14
9.1	Core conclusion	14
9.2	Economic intuition	14
9.3	Incremental contribution revisited	14
9.4	Implications for innovation policy and organization	15
9.5	Limitations	15
9.6	Takeaway	15

1 Big picture

- **Question.** Is innovation becoming harder because each generation of innovators must first master a larger stock of prior knowledge? The paper asks how an expanding knowledge frontier changes innovator education, specialization, teamwork, and long-run growth.
- **Core idea.** Knowledge accumulation has a hidden cost. To stand on the shoulders of giants, innovators must first climb up those giants' backs. As the knowledge stock rises, the educational burden of reaching the frontier rises as well.
- **Two individual responses.** Innovators can respond by learning more or by narrowing their expertise. Learning more delays active innovation; narrowing expertise reduces individual scope and increases the need for teams.
- **Empirical setting.** The paper uses the Hall, Jaffe, and Trajtenberg patent database, covering all USPTO utility patents from 1963 to 1999, and augments it with inventor age data for a large subset of U.S. inventors.
- **Three main outcomes.** The paper studies team size, age at first innovation, and specialization. Team size is the number of inventors listed on a patent. Age at first innovation proxies for education or preparation time. Specialization is proxied by the probability that a solo inventor switches technological fields between consecutive patents.
- **Six facts.** Over time, age at first innovation rises, specialization rises, and team size rises. In cross-section, team size and specialization are greater in deeper fields, but age at first innovation varies surprisingly little across fields.
- **Theory.** The model organizes knowledge into application areas and foundational knowledge circles. Innovators choose an application area, a position in the knowledge space, and a breadth of expertise. The burden of knowledge is the cost of reaching the frontier in that area.
- **Main mechanism.** In equilibrium, innovators arbitrage across application areas so that life-time income is equalized. This makes total education similar across fields, while deeper fields induce narrower expertise and larger teams.
- **Growth implication.** The burden of knowledge acts like a micro-foundation for “ideas are getting harder to find.” It can explain why R&D effort increased dramatically in the 20th century without a corresponding acceleration in TFP growth.
- **Economic takeaway.** Innovation becomes increasingly collective and specialized not merely because teams are fashionable, but because the frontier itself becomes harder for any single person to cover.

2 Incremental contribution of the paper

- **From aggregate innovation difficulty to individual mechanisms.** Existing growth models often use reduced-form fishing-out or difficulty parameters. This paper provides a micro mechanism: knowledge accumulation raises the human-capital burden faced by innovators.
- **A new set of micro facts.** The paper documents time-series and cross-sectional facts about inventor age, specialization, and teamwork. These facts connect individual careers to aggregate growth dynamics.

- **Augmented patent data.** The paper builds detailed individual patent histories and adds age information for more than 55,000 inventors using name and zip-code matching. This allows an outcome-based measure of the transition into innovative activity.
- **A direct measure of knowledge depth.** The paper uses the size of the backward citation tree behind a patent as a proxy for the amount of underlying knowledge in a technological area.
- **Specialization as field-switching.** The paper operationalizes specialization using whether solo inventors switch technological categories across consecutive patents. This is a clever proxy because technological fields are assigned to patents rather than people.
- **Unified explanation of six facts.** The model jointly explains rising education, rising specialization, rising teamwork, greater specialization and teamwork in deeper fields, and the cross-sectional stability of age at first innovation.
- **No-scale-effects contribution.** The model explains how the economy can devote exponentially more researchers to R&D while productivity growth remains flat. Increasing effort is needed just to offset rising knowledge burdens.
- **Organization-of-innovation contribution.** The paper links the burden of knowledge to the decline of the “Renaissance Man,” the rise of teams, and the narrowing of individual expertise.
- **Policy implication.** The model implies that the efficiency of education and intergenerational knowledge transfer is a first-order determinant of innovation organization and long-run growth.

3 Data and measurement

3.1 Patent data

- The core data set is Hall, Jaffe, and Trajtenberg’s USPTO patent data set.
- The full data contain every utility patent issued by the USPTO from 1963 to 1999.
- The 1963–1999 universe includes about 2.9 million patents. The post-1975 part, where citation and inventor-name data are available, includes about 2.1 million patents.
- Patent information includes grant date, application year, technological category, citations, and inventor names and addresses for patents granted after 1975.
- Technological fields are measured at three levels: the USPTO 414-class definition, Hall et al.’s 36-category classification, and Hall et al.’s broad 6-category classification.

Interpretation: The data allow the paper to study both time trends and cross-field variation in innovative activity. The main strength is scale; the main limitation is that patents are an imperfect measure of innovation and do not cover all innovative work.

3.2 Inventor histories and age data

- Inventors are identified by last name, first name, and middle initial.
- The paper builds patent histories for each inventor and observes consecutive patent applications.

- The data do not contain birth dates, so the author matches inventor names and zip codes to public records from AnyBirthday.com.
- Of 224,152 inventors with a U.S. zip code, 56,281 have a unique birth-date match.
- The age subset is not random, so the paper discusses selection and uses robustness checks. The age results should be read with this limitation in mind.

Interpretation: The age data are valuable because they allow a direct outcome-based measure of when inventors first become active. The trade-off is potential selection from public-record matching.

3.3 Main outcome variables

- **Team size.** The number of inventors listed on a patent.
- **Age at first innovation.** The age at which an inventor first appears in the patent data. The paper focuses on post-1985 first appearances and age windows such as 25–35 or 23–33 to reduce left-censoring problems.
- **Field jump.** The probability that a solo inventor switches technological fields between two consecutive patents. A lower field-jump probability indicates greater specialization.
- **Time lag.** The delay between consecutive applications from the same inventor.
- **Tree size.** The logarithm of the number of nodes in a patent’s backward citation tree. This is the main proxy for knowledge depth.

Interpretation: The paper’s measurement strategy turns abstract concepts - education, specialization, teamwork, and knowledge depth - into observable patent-based variables.

3.4 Why solo inventors matter for specialization

- Technological categories are assigned to patents, not to people.
- For team patents, a field switch could reflect a different teammate rather than a change in an individual’s expertise.
- The paper therefore focuses on solo inventors when using field jumps as a specialization measure.

Interpretation: This sample restriction makes the specialization proxy cleaner, but it also means the field-jump results apply most directly to solo inventors.

4 Empirical specifications and evidence

4.1 Time-series regressions

The time-series regressions relate each outcome to application year. A generic specification is

$$Y_{pt} = \alpha + \beta \text{ApplicationYear}_{pt} + \Gamma' X_{pt} + \varepsilon_{pt},$$

where Y_{pt} is team size, age at first innovation, or field-jump status, and X_{pt} includes controls such as technological category, assignee code, foreign-source indicators, team size, or time lag depending on the outcome.

- **Team size.** Table 2 shows a strong upward trend in inventors per patent. The basic coefficient is about 0.029 inventors per year, or roughly 14% per decade relative to the period mean.
- **Age at first innovation.** Table 3 shows an upward age trend of roughly 0.066 years per application year, or about two-thirds of a year per decade.
- **Specialization.** Table 4 shows a declining probability of switching fields. Using the 414-category field definition, the probability of field jump falls by about 0.003 per year, or about 6% per decade relative to the mean.
- **Robustness.** The trends survive broad and narrow technological controls, domestic-versus-foreign splits, assignee-code controls, and alternative age windows.

Interpretation: The time-series evidence is consistent with the burden-of-knowledge idea: newer cohorts enter innovation later, work in larger teams, and appear more specialized.

4.2 Cross-sectional regressions using knowledge depth

For cross-sectional tests, the paper uses normalized variation in citation-tree size:

$$\text{TreeSize}_{pt} = \frac{\log(\text{nodes in citation tree}_{pt}) - \overline{\log(\text{nodes})}_t}{sd_t(\log(\text{nodes}))}.$$

A generic specification is

$$Y_{pt} = \alpha_t + \beta_1 \text{TreeSize}_{pt} + \beta_2 \text{TreeSize}_{pt}^2 + \Gamma' X_{pt} + \varepsilon_{pt}.$$

- **Teamwork.** Table 6 shows that larger citation trees are associated with more inventors per patent. The relationship is positive, robust, and convex.
- **Age.** Table 7 shows little robust relationship between tree size and age at first innovation. This is the surprising cross-sectional fact that motivates the model.
- **Specialization.** Table 8 shows that solo inventors are less likely to switch fields when their initial patent has a larger citation tree. Deeper knowledge is associated with greater specialization.
- **Controls.** The regressions include application-year dummies, technological field controls, foreign-source indicators, direct-citation controls, and time-between-application controls depending on the test.

Interpretation: Deeper fields use more teams and narrower individual expertise, but not visibly longer pre-innovation education. The model explains this by endogenous sorting across application areas.

5 The model

5.1 Knowledge space

The model organizes knowledge into application areas, such as airplanes or drugs. Each application area has a circle of foundational knowledge. An innovator chooses

$$\omega = (j, s_j, b_j),$$

where j is the application area, s_j is the position in that area's knowledge space, and b_j is breadth of expertise.

The education cost is

$$E(\omega, \tau) = (b_j D_j(\tau))^\epsilon,$$

where $D_j(\tau)$ is the depth of knowledge in area j . Knowledge depth evolves as

$$D_j(t) = \hat{D}_j X(t)^\delta.$$

Interpretation: Broader expertise and deeper knowledge both increase the cost of becoming an innovator.

5.2 Idea arrival and crowding

Once educated, innovators receive ideas at rate

$$\lambda(\omega, t) = \hat{A}_j X(t)^\chi L_j(t, s_j)^{-\sigma} b_j^\beta.$$

- $\hat{A}_j$ captures application-specific technological opportunity.
- $X(t)^\chi$ captures whether existing knowledge makes future innovation easier or harder.
- $L_j(t, s_j)^{-\sigma}$ captures crowding or duplication among similar innovators.
- b_j^β captures the productivity benefit of broader expertise.

Interpretation: Innovators value breadth because it raises innovative potential, but breadth is costly when the knowledge frontier is deep.

5.3 Team formation

The model assumes that implementing an idea requires covering the foundational knowledge circle for that application area. If a single innovator's breadth does not cover the whole circle, a team must be formed. Thus narrower expertise naturally implies larger teams.

For an equilibrium breadth $b_j^*(\tau)$, team size is approximately

$$team_j(\tau) = \begin{cases} 1, & b_j^*(\tau) > 1, \\ \lceil 1/b_j^*(\tau) \rceil, & b_j^*(\tau) \leq 1. \end{cases}$$

Interpretation: Teamwork is not an auxiliary assumption. It is the organizational consequence of specialization.

5.4 Equilibrium implications

Along a balanced growth path, the model produces several key results.

- **Education equalization across fields.** Innovators in the same cohort choose the same total education expenditure across application areas, even if some fields are deeper or have better opportunities.
- **Rising education over time.** Total education grows across cohorts at the productivity growth rate.
- **Narrower breadth in deeper fields.** Since education is equalized, fields with greater knowledge depth induce narrower breadth.
- **Increasing specialization over time.** Breadth falls over time if the knowledge depth rises fast enough, specifically if $\delta > 1/\epsilon$.
- **Larger teams in deeper fields.** Team size is larger when breadth is narrower, so deeper fields use larger teams.

Interpretation: The model matches the main puzzle: education is similar across fields because innovators arbitrage across fields, while breadth and teamwork adjust to field-specific knowledge depth.

5.5 Growth equation

The steady-state productivity growth rate is

$$g = \frac{1 - \sigma}{1 - \chi + \beta \left(\delta - \frac{1}{\epsilon} \right)} g_L.$$

- g_L is population growth.
- χ captures whether existing knowledge makes future innovation easier or harder.
- $\beta(\delta - 1/\epsilon)$ captures the burden-of-knowledge channel.
- If the knowledge burden rises, specialization increases and the effective contribution of each innovator can fall.

Interpretation: The burden of knowledge works like a micro-foundation for fishing-out effects. Even if more researchers enter R&D, individual innovative capacity can decline enough to eliminate scale effects.

6 Identification and empirical design

6.1 What the empirical design can identify

- The empirical evidence identifies strong descriptive patterns in inventor behavior across time and fields.

- The paper does not claim randomized causal identification. The evidence is best read as systematic facts that a unified model explains.
- The model provides a coherent interpretation for patterns that are otherwise difficult to reconcile jointly.

6.2 Key empirical challenges

- **Left-censoring of first innovations.** Because inventor histories begin in 1975, an inventor's first observed patent might not be the true first patent. The paper therefore focuses on later first appearances and age windows.
- **Name matching.** Inventors are identified by names and initials, which can create false merges or false splits.
- **Age-data selection.** Birth-date matching is available only for a selected subset of U.S. inventors with zip-code information.
- **Citation-tree proxy.** Tree size is an imperfect measure of knowledge depth because citation practices and data truncation vary across time and fields.
- **Specialization proxy.** Field jumps are a noisy measure of expertise, especially for team patents. The paper therefore focuses on solo inventors.

6.3 Why the evidence is still useful

- The main trends are large and robust to many controls.
- The same qualitative patterns appear across broad and narrow technological categories.
- The cross-sectional facts use within-year normalized tree-size variation, reducing the effect of time trends in citations.
- Appendix checks show that the age subset is selected, but many differences are explainable with observables.

Interpretation: The empirical strategy is strongest as a fact-building exercise. The paper's contribution is to show that these facts line up with a disciplined equilibrium mechanism.

7 Tables and figures

7.1 Table 1 and Figure 1: motivating age trends and basic patent trends

Read:

- Table 1 shows that Nobel Prize winners' age at great achievement rises by about 5.83 years per century, and age at highest degree rises by about 4.11 years per century.
- Figure 1 shows three basic trends in the patent data: age at first innovation rises, the probability of field jump for fast innovators falls, and average inventors per patent rises.

- Together, the two visuals motivate the paper’s central empirical claim: innovators are entering later, specializing more, and working in larger teams.

Economic message: The burden of knowledge is visible both in elite scientific careers and in patenting behavior. The figures suggest that innovation increasingly requires longer preparation and more division of labor.

TABLE 1
Age trends among Nobel Prize winners

	Dependent variable	
	Age at great achievement	Age at highest degree
Year of great achievement (in 100's)	5.83*** (1.37)	—
Year of highest degree (in 100's)	—	4.11*** (0.61)
Number of observations	544	505
Time span	1873–1998	1858–1990
Average age	38.6	26.5
R ²	0.027	0.084

Notes: (i) This table borrows from Jones (2005a). Age trends are measured in years per century. S.E. are given in parentheses. (ii) Nobel Prize winners include all winners in Physics, Chemistry, Medicine, and Economics. Age at great achievement is age when contribution is made (not later age when prize is awarded).
***Indicates 99% confidence level.

(a) Table 1: Age trends among Nobel Prize winners




FIGURE 1
Basic time trends

(b) Figure 1: Basic time trends

7.2 Figure 2 and Table 2: data window and teamwork trend

Read:

- Figure 2 summarizes the patent-data window: basic patent data are available from 1963, while citations and inventor names are available after 1975.
- Table 2 estimates the trend in inventors per patent. The upward trend is large, robust to technological controls, assignee controls, and domestic-versus-foreign splits.
- The baseline estimate implies about 0.029 additional inventors per patent per year.

Economic message: Teamwork rises broadly, not just in one field or ownership category. This supports the interpretation that the organization of innovation is changing economy-wide.




FIGURE 2
Summary of available data

(a) Figure 2: Summary of available data

TABLE 2
Trends in inventors per patent

	Dependent variable: inventors per patent						
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Application year	0.0293 (0.0001)	0.0261 (0.0001)	0.0262 (0.0001)	0.0251 (0.0001)	0.0244 (0.0002)	0.0306 (0.0002)	0.0180 (0.0003)
Foreign patent	—	0.444 (0.002)	0.416 (0.002)	0.141 (0.004)	0.146 (0.004)	U.S. only	Foreign only
Technological field controls	—	Yes	—	—	—	—	—
Broad	—	—	—	—	—	—	—
Narrow	—	—	—	—	—	—	—
Assignee code	—	Yes	Yes	Yes	Yes	Yes	Yes
Number of observations	2,016,377	2,016,377	2,016,377	2,016,377	1,506,956	1,123,310	893,067
Period	1975–1999	1975–1999	1975–1999	1975–1999	1975–1996	1975–1999	1975–1999
Mean of dependent variable	2.03	2.03	2.03	2.03	1.97	1.82	2.29
Per decade trend as % of period mean	14.4	12.9	12.9	12.4	12.4	16.8	7.9
R ²	0.02	0.08	0.10	0.12	0.13	0.12	0.10

Note: (i) Regressions are ordinary least squares with S.E. in parentheses. Specifications (1) through (4) consider the entire universe of patents applied for between 1975 and 1999. Specification (5) considers only patents that were granted within 3 years after application (see discussion in text). Specifications (6) and (7) present separate trends for domestic and foreign source patents. (ii) Foreign patent is a dummy variable to indicate whether the first inventor listed with the patent has an address outside the U.S. (iii) “Broad” technological controls include dummies for each of the six categories in the most aggregated technological classification of Hall *et al.* “Narrow” technological controls include dummies for each category of their 36-category classification. (iv) Upward trends persist when run separately for each technological field. Using the broad classification (six categories), the trends range from a low of 0.013 for “Other” to a high of 0.037 for “Chemical.” Using the narrower classification scheme (36 categories), the trends range from a low of 0.007 for “Apparel & Textile” to 0.051 for “Organic Compounds.” The smallest *t*-statistic for any of these trends is 7.76. (v) Assignee code controls are seven dummy variables that define who holds the rights to the patent. Most patent rights are held by U.S. or foreign corporations (80%), while a minority remain unassigned (17%) at the time the patent is issued. Table B.2 describes the assignee codes in further detail. Running the time trends separately for the individual assignee codes shows that the team size trends range from a low of 0.005 for the unassigned category to a high of 0.039 for U.S. non-government institutions. The lowest *t*-statistic for any of these trends is 5.38.

(b) Table 2: Trends in inventors per patent

7.3 Tables 3 and 4: age at first innovation and specialization over time

Read:

- Table 3 shows that age at first innovation rises by roughly 0.066 years per application year, or about two-thirds of a year per decade.
- The age trend is robust to technological controls, assignee-code controls, team size, and alternative age windows.
- Table 4 shows that the probability of switching technological fields declines over time. This indicates rising specialization.
- The field-jump result is robust across 414-, 36-, and 6-category field definitions.

Economic message: Successive cohorts spend longer before entering innovation and then operate in narrower technological domains. This is the behavioral signature of a rising knowledge burden.

TABLE 3
Trends in age at first innovation

	Dependent variable: age at application						
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Application year	0.0657 (0.0095)	0.0666 (0.0095)	0.0671 (0.0095)	0.0671 (0.0099)	0.0687 (0.0097)	0.0530 (0.0107)	0.0584 (0.0109)
Technological field controls	—	Yes	—	—	—	—	—
Broad	—	—	—	—	—	—	—
Narrow	—	—	—	—	—	—	—
Assignee code	—	Yes	Yes	Yes	Yes	Yes	Yes
Team size	—	—	—	—	—	—	—
Number of observations	6541	6541	6541	6541	5102	5102	5102
Period	1985–1999	1985–1999	1985–1999	1985–1999	1985–1999	1985–1999	1985–1999
Age range	25–35	25–35	25–35	25–35	25–35	25–35	25–35
Mean of dependent variable	31.0	31.0	31.0	31.0	29.3	29.3	29.3
Per decade trend as % of period mean	2.1	2.1	2.2	2.2	1.8	2.0	2.0
R ²	0.007	0.010	0.020	0.020	0.021	0.005	0.018

Note: (i) Regressions are ordinary least squares, with S.E. in parentheses. All regressions observe only at those innovators for whom we have age data and who appear for the first time in the data set in or after 1985. Specifications (1) through (3) consider those innovators who appear for the first time between ages 25 and 35. Specifications (4) and (7) consider those innovators who appear for the first time between ages 23 and 33. (ii) Broad technological controls include dummies for each of the six categories in the most aggregated technological classification Hall *et al.* Here, 12 categories show significant upward trends. Sample sizes drop considerably when the data are divided into these 36 categories. The one case of a significant downward trend (category no. 25, Computer Peripherals) has 42 observations. (iii) Assignee code controls are seven dummy variables that define who holds the rights to the patent. Table B.2 describes the assignee codes in further detail. The upward age trends persist when run separately for each assignee code and are similar in magnitude to the trends in Table 2.

(a) Table 3: Trends in age at first innovation

TABLE 4
Trends in probability of field jump

	Dependent variable: probability of switching technological field							
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Application year	−3.4e−3 (0.19e−3)	−3.2e−3 (0.19e−3)	−2.5e−3 (0.19e−3)	−2.8e−3 (0.17e−3)	−3.1e−3 (0.17e−3)	−3.1e−3 (0.12e−3)	−3.1e−3 (0.11e−3)	−3.0e−3 (0.11e−3)
Foreign patent	—	0.0076 (0.0039)	—	−0.0041 (0.0038)	—	0.0002 (0.0035)	—	−0.0005 (0.0029)
Time between applications	—	0.0225 (0.0012)	—	0.0206 (0.0012)	—	0.0154 (0.0011)	—	0.0228 (0.0044)
Technological field controls (first patent)	—	—	—	—	—	—	—	—
Assignee code (first patent)	—	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Number of observations	215,855	215,855	215,855	215,855	215,855	359,405	359,405	359,405
Period	1975–1993	1975–1993	1975–1993	1975–1993	1975–1993	1975–1999	1975–1999	1975–1999
Mean of dependent variable	0.535	0.535	0.423	0.423	0.294	0.294	0.556	0.556
Per decade trend as % of period mean	−6.4	−6.0	−5.9	−6.4	−6.5	−7.8	−9.4	−5.6
Pseudo-R ²	0.0011	0.018	0.0006	0.019	0.0005	0.017	0.004	0.026

Note: (i) Results are for probit estimation, with coefficients reported at mean values and S.E. in parentheses. The coefficient for the foreign dummy is reported over the 0–1 range. (ii) The dependent variable is 0 if an inventor does not switch fields between two consecutive innovations. The dependent variable is 1 if the inventor does switch fields. Column headings define the field classification used to determine the dependent variable. “114” indicates the 414-category technological class definition of the U.S. Patent and Trademark Office; “36” and “6” refer to the aggregated measures defined by Hall *et al.* (2001). (iii) Specifications (1) through (6) consider “first” innovators—only those consecutive patents with no more than 3 years between applications and with no more than 3 years delay between application and grant (see discussion in text). Specifications (7) and (8) consider all consecutive patents. (iv) Technological field controls are dummies for the 36 categories defined by Hall *et al.* (2001). The reported regressions use the technological field of the initial patent. Using the field of the second patent has no effect on the results. Running the regressions separately by technology category shows that the trends persist in six of the six categories using the broad technology classification of Hall *et al.* and 34 of the 36 categories using the narrow classification of Hall *et al.*, with significant trends in 20. (v) Assignee code controls are seven dummy variables that define who holds the rights to the patent. Table B.2 describes the assignee codes in further detail. The declining probability of field jump persists when the trend is examined within each assignee code, although the significance of the trend disappears in the rarer classifications.

(b) Table 4: Trends in probability of field jump

7.4 Tables 5 and 6: cross-field patterns and tree-size evidence for teamwork

Read:

- Table 5 compares technological categories. Age at first innovation is strikingly similar across fields, while inventors per patent and field-jump probabilities vary substantially.

- The similarity in age across fields is the surprising fact that a simple burden-of-knowledge story alone would not predict.
- Table 6 shows that patents with larger citation trees have more inventors. The positive relationship remains with year dummies, field controls, domestic/foreign splits, and direct-citation controls.

Economic message: Deeper areas of knowledge use more teamwork, but they do not necessarily show higher age at first innovation. The model explains this by allowing innovators to sort across application areas while adjusting breadth.

TABLE 5
Mean differences across technological categories

	Technological classification (Hall <i>et al.</i> , 2001)	Code	Age at first innovation		Inventors per patent		Probability of field jump	
			Obs.	Mean	Obs.	Mean	Obs.	Mean
6	36							
Chemical (1)	Agriculture, Food, Textiles	11	12	31.1	18,351	2.45	2461	0.48
	Coating	12	52	29.2	32,820	2.24	4336	0.64
	Gas	13	17	29.9	10,047	1.96	1692	0.59
	Organic Compounds	14	49	29.4	78,188	2.57	8010	0.32
	Resins	15	43	29.4	74,993	2.52	7695	0.37
	Miscellaneous—Chemical	19	329	29.3	214,854	2.23	29,721	0.43
Entire category			502	29.4	429,253	2.35	53,915	0.43
Computers & Communications (2)	Communications	21	270	29.2	98,046	1.99	15,107	0.41
	Computer Hardware & Software	22	169	29.8	83,094	2.26	10,259	0.44
	Computer Peripherals	23	38	29.2	22,809	2.37	2772	0.51
	Information Storage	24	45	29.0	43,182	2.21	6778	0.40
	Entire category		522	29.4	247,131	2.15	34,916	0.42
Drugs & Medical (3)	Drugs	31	74	29.9	77,210	2.91	7181	0.25
	Surgery & Medical Instruments	32	276	29.8	62,192	1.86	12,385	0.29
	Biotechnology	33	48	30.5	29,638	2.75	2223	0.36
	Misc—Drugs & Medical	39	71	29.2	14,356	1.66	3488	0.35
	Entire category		469	29.8	183,396	2.43	25,277	0.29
Electrical & Electronic (4)	Electrical Devices	41	116	29.2	65,500	1.77	12,817	0.48
	Electrical Lighting	42	88	29.6	33,769	1.97	5739	0.43
	Measuring & Testing	43	117	29.1	62,021	1.94	10,083	0.51
	Nuclear & X-rays	44	56	29.5	32,402	2.08	4681	0.50
	Power Systems	45	124	29.4	73,849	1.94	13,086	0.51
	Semiconductor Devices	46	51	29.3	47,123	2.25	7207	0.34
	Misc—Electrical	49	103	29.1	52,206	1.97	9004	0.51
Entire category			655	29.3	366,870	1.97	62,617	0.48
Mechanical (5)	Materials Processing & Handling	51	243	29.4	108,873	1.79	21,821	0.48
	Metal Working	52	89	28.8	63,669	2.12	10,454	0.54
	Motors, Engines & Parts	53	86	29.4	78,585	1.85	16,221	0.41
	Optics	54	57	29.0	51,102	2.15	8159	0.37
	Transportation	55	279	28.9	61,501	1.66	12,004	0.45
	Misc—Mechanical	59	458	29.1	103,855	1.64	22,513	0.49
Entire category			1212	29.1	467,585	1.84	91,172	0.46
Others (6)	Agriculture, Husbandry, Food	61	248	29.1	44,718	1.76	7644	0.40
	Amusement Devices	62	267	29.5	22,227	1.41	4273	0.37
	Apparel & Textile	63	204	29.2	35,001	1.57	7616	0.37
	Earth Working & Wells	64	98	29.7	29,645	1.69	6599	0.36
	Furniture, House Fixtures	65	352	29.1	43,499	1.42	9416	0.50
	Heating	66	55	30.0	28,267	1.76	6065	0.48
	Pipes & Joints	67	46	29.2	18,444	1.58	4448	0.61
	Receptacles	68	297	29.4	43,353	1.51	10,105	0.47
	Misc—Others	69	848	29.2	179,925	1.74	35,342	0.48
	Entire category		2415	29.3	445,079	1.65	91,508	0.46

Notes: (i) Age at first innovation includes observations of those innovators who appear after 1985 in the data set and between the ages of 23 and 33. Results are similar, with higher mean and even less variance, for 25- to 35-year-olds. (ii) Probability of field jump is probability of switching categories for solo innovators using 36-category measure. Obs., observation.

(a) Table 5: Mean differences across technological categories

TABLE 6 Inventors per patent versus tree size						
	Dependent variable: inventors per patent					
	(1)	(2)	(3)	(4)	(5)	(6)
Normalized variation in tree size	0.0849 (0.0010)	0.0961 (0.0010)	0.0995 (0.0011)	0.120 (0.001)	0.133 (0.001)	0.107 (0.001)
Normalized variation in tree size, squared	0.0069 (0.0007)	0.0545 (0.0007)	0.0545 (0.0007)	0.0341 (0.0007)	0.0257 (0.0009)	0.0356 (0.0011)
Foreign patent	—	0.446 (0.002)	0.442 (0.002)	0.420 (0.002)	U.S. only	Foreign only
Normalized variation in direct citations made	—	—	-0.0394 (0.0011)	—	—	—
Technological field controls	—	—	—	Yes	Yes	Yes
Application year dummies	Yes	Yes	Yes	Yes	Yes	Yes
Number of observations	1,969,908	1,969,908	1,969,908	1,969,908	1,103,402	866,506
Period	1975-1999	1975-1999	1975-1999	1975-1999	1975-1999	1975-1999
Mean of dependent variable	2.02	2.02	2.02	2.02	1.82	2.27
R ²	0.026	0.050	0.050	0.100	0.090	0.083

Notes: (i) Regressions are ordinary least squares with S.E. in parentheses. Specifications (1) through (4) consider the entire universe of patents applied for between 1975 and 1999. Specifications (5) and (6) consider separately patents from domestic versus foreign sources. Specification (7) considers cross-sections from the later part of the time period. (ii) Normalized variation in tree size is the deviation from the year mean tree size, divided by the year S.D. in tree size. "Tree size" is the log of the number of nodes in the citations tree behind any patent. (iii) Normalized variation in direct citations made captures variation in the number of citations to prior art listed on a patent application. It is the deviation from the year mean number of citations, divided by the year S.E. in the number of citations. (iv) Technological field controls include dummies for each of 36-category measure of Hall *et al.* (v) The number of observations here is slightly smaller than for the time trend analysis in Table 2 because a few patents do not cite other U.S. patents; hence, no citation tree can be built; these patents are dropped from the analysis.

(b) Table 6: Inventors per patent versus tree size

7.5 Tables 7 and 8: tree size, age, and specialization

Read:

- Table 7 shows that the relationship between citation-tree size and age at first innovation is weak and unstable.
- This is the key cross-sectional puzzle: deeper fields do not clearly require later first invention.
- Table 8 shows that citation-tree size negatively predicts field jumping by solo inventors. Deeper fields are associated with lower switching and greater specialization.

Economic message: The cross-sectional evidence says that knowledge depth is absorbed through narrower breadth and larger teams rather than through substantially more total education within a cohort.

TABLE 7
Age versus tree size

	Dependent variable: age at application for first patent							
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Normalized variation in tree size	-0.007 (0.032)	-0.005 (0.036)	0.114 (0.035)	0.084 (0.040)	0.099 (0.043)	0.097 (0.039)	0.113 (0.046)	0.030 (0.026)
Team size	—	-0.054 (0.027)	—	-0.036 (0.030)	-0.038 (0.030)	-0.024 (0.025)	0.008 (0.035)	-0.029 (0.039)
Normalized variation in direct citations made	—	—	—	0.064 (0.044)	—	—	—	—
Technological field controls	—	Yes	—	Yes	—	Yes	Yes	Yes
Application year dummies	Yes	Yes	—	Yes	Yes	Yes	Yes	Yes
Number of observations	6486	6446	8058	8058	8058	8434	3630	3588
Period	1985-1999	1985-1999	1985-1999	1985-1999	1985-1999	1975-1999	1985-1999	1985-1999
Age range	25-35	25-35	23-33	23-33	23-33	23-33	21-31	28-33
Mean of dependent variable	31.0	31.0	29.34	29.3	29.2	29.2	27.7	30.7
R ²	0.009	0.022	0.009	0.021	0.012	0.020	0.025	0.020

Notes: (i) Regressions are ordinary least squares, with S.E. in parentheses. All regressions observe only at those innovators for whom we have age data. Specifications (1) and (2) consider first innovations in the 25-35 age window. Specifications (3) through (6) consider innovators in the 23-33 age window. Specification (7) considers slightly younger innovators, and Specification (8) considers the latter half of the 23-33 age window. Specification (6) considers cross-sections pooled over the entire time period; the other specifications focus on the post-1985 period, for which we can be confident that we are witnessing an innovator's first patent. (ii) Normalized variation in tree size is the deviation from the year mean tree size, divided by the year S.E. in tree size. Tree size is the log of the number of nodes in the citations tree behind any patent. (iii) Normalized variation in direct citations made captures variation in the number of citations to prior art listed on a patent application. It is the deviation from the year mean number of citations, divided by the year S.E. in the number of citations. (iv) The number of observations here is slightly smaller than for the time trend analysis in Table 3 because a few patents do not cite other U.S. patents; hence, no citation tree can be built; these patents are dropped from the analysis. (v) Technological field controls include dummies for each of 36-category measure of Hall *et al.*

(a) Table 7: Age versus tree size

TABLE 8
Field jump versus tree size

	Dependent variable: probability of switching technological field					
	(1)	(2)	(3)	(4)	(5)	(6)
Normalized variation in tree size	-0.0072 (0.0008)	-0.0074 (0.0008)	-0.0059 (0.0008)	-0.0095 (0.0009)	-0.0144 (0.0012)	-0.0184 (0.0017)
Foreign patent	—	—	-0.0125 (0.0018)	-0.0108 (0.0018)	-0.0129 (0.0023)	-0.0135 (0.0032)
Time between applications	—	—	0.0226 (0.0004)	0.0232 (0.0004)	0.0215 (0.0012)	0.0143 (0.0017)
Technological field controls (first patent)	—	—	—	Yes	Yes	Yes
Application year dummies	Yes	Yes	Yes	Yes	Yes	Yes
Number of observations	353,762	353,762	353,762	353,762	212,274	110,511
Period	1975-1999	1975-1999	1975-1999	1975-1999	1975-1993	1985-1993
Mean of dependent variable	0.551	0.551	0.551	0.551	0.536	0.520
Pseudo-R ²	0.0039	0.0039	0.0117	0.0251	0.0171	0.0159

Notes: (i) Results are for probit estimation, with coefficients reported at mean values and S.E. in parentheses. The coefficient for the foreign dummy is reported over the 0-1 range. Only solo inventors are considered. Specifications (1) through (4) consider the entire set of solo inventors. Specification (5) considers only those solo inventors who meet the criteria in Specifications (1) through (6) in Table 4 (to help control for any truncation bias in the specialization measure—see the discussion of Table 4 in the text). Specification (6) considers the same data as Specification (5) but only examines cross-sections in the later part of the time period. (ii) The dependent variable is 0 if an inventor does not switch fields between two consecutive innovations. The field is defined using the 414-category technological class definition of the U.S. Patent and Trademark Office. (iii) Normalized variation in tree size is the deviation from the year mean tree size, divided by the year S.E. in tree size. Tree size is the log of the number of nodes in the citations tree behind any patent. (iv) Technological field controls include dummies for each of 36-category measure of Hall *et al.*

(b) Table 8: Field jump versus tree size

7.6 Figure 3 and Table B.1: model geometry and age-data selection

Read:

- Figure 3 illustrates the model's knowledge space. The horizontal circle represents the range of foundational knowledge types, and the vertical dimension represents depth $D_j(t)$.
- An innovator chooses a starting point s_j and breadth b_j . Greater depth makes a given breadth more expensive to acquire.
- Table B.1 shows the stages of selection into the age-data sample. Only a subset of U.S. inventors with zip codes can be uniquely matched to birth-date records.

Economic message: The figure is the model's core intuition: as the knowledge cylinder gets taller, covering a fixed arc requires more education. The appendix table reminds the reader that the age results depend on a selected but analyzable subset.




FIGURE 3

A circle of knowledge

(a) Figure 3: A circle of knowledge

TABLE B.1

Number of observations at each stage of selection

	Number of observations	Percentage of row (3)	Percentage of row (4)	Percentage of row above
(1) Patents granted	2,139,313			
(2) Inventors worldwide	4,301,229			
(3) Unique inventors worldwide	1,411,842			
(4) Unique inventors with U.S. address	752,163	53.3		53.3
(5) Unique inventors with U.S. address, zip code	224,152	15.9	29.8	29.8
(6) Unique inventors with U.S. address, zip code, unique match from http://www.AnyBirthday.com	56,281	4.0	7.5	25.1

Notes: (i) Observation counts consider the 1975–1999 period. (ii) A “unique inventor” is defined by having same first name, last name, and middle initial.

(b) Table B.1: Number of observations at each stage of selection

7.7 Tables B.2 and B.3: appendix checks on the age sample

Read:

- Table B.2 compares patent-right assignment across samples and shows that zip-code and birth-date availability are related to assignment status.
- Table B.3 compares inventors per patent across sample groups. Selection into the age sample is visible, but the differences shrink when controls are added.
- These appendix tables support the paper’s claim that the age subset is not random, but the main age patterns are not mechanically driven by obvious sample composition alone.

Economic message: The paper is transparent about measurement limits. The age data are not a clean random sample, so the author checks whether selection is large enough to overturn the interpretation.

TABLE B.2

The assignment of patent rights

Assignment status	All patents	U.S. patents	U.S. patents no zip code	U.S. patents zip code	Birth data	
					Direct match	Other patents
Unassigned (%)	17.2	22.4	0.4	98.3	97.9	26.6
U.S. non-government organization (%)	43.9	72.9	94.1	0.0	0.0	65.7
Non-U.S. non-government organization (%)	36.2	1.1	1.4	0.0	0.0	3.4
Other assignment (%)	2.7	3.5	4.1	1.7	2.1	4.4

Notes: (i) The first column considers all patent observations in the 1975–1999 period (2.1 million observations). (ii) U.S. patents are those for which first inventor listed with the patent has a U.S. address. (iii) The birth data columns consider those U.S. patents with zip code information for which <http://www.AnyBirthday.com> produced a birth date. The first birth data column considers the specific patents on which <http://www.AnyBirthday.com> was able to match. The last column considers all other patents by that innovator, identifying the innovator by last name, first name, and middle initial. (iv) Unassigned patents are those for which the patent rights were still held by the original inventor(s) at the time the patent was granted; these patents may or may not have been assigned after the grant date. (v) Non-government organizations are mainly corporations but also include universities. (vi) Other assignment includes assignments to: (a) U.S. individuals, (b) non-U.S. individuals, (c) the U.S. government, and (d) non-U.S. governments.

(a) Table B.2: Assignment of patent rights

TABLE B.3

Inventors per patent, mean differences between samples

	Dependent variable: inventors per patent				
	(1)	(2)	(3)	(4)	(5)
U.S. address dummy	−0.315 (0.0020)	−0.339 (0.0020)	−0.300 (0.0020)	−0.124 (0.0049)	−0.103 (0.0048)
U.S. address and zip code dummy	−0.786 (0.0033)	−0.670 (0.0033)	−0.769 (0.0032)	−0.155 (0.0069)	−0.176 (0.0066)
U.S. address, zip code, and http://www.AnyBirthday.com direct match dummy	0.237 (0.0068)	0.246 (0.0067)	0.212 (0.0067)	0.243 (0.0067)	0.228 (0.0066)
Constant	2.28 (0.0014)	2.57 (0.0023)	1.96 (0.0052)	1.45 (0.0042)	1.56 (0.0067)
Technological category dummies	No	Yes	No	No	Yes
Grant year dummies	No	No	Yes	No	Yes
Assignee code dummies	No	No	No	Yes	Yes
R ²	0.0555	0.0825	0.0756	0.0757	0.1162

Notes: (i) Regressions consider means in the entire data set (2.1 million patent observations), covering the 1975–1999 time period. S.E. are given in parentheses. (ii) Dummy variables are nested: the second row captures a subset of the first. The third row captures a subset of the second. (iii) Innovators for whom <http://www.AnyBirthday.com> produces a birth date are often involved with multiple innovations over the 1975–1999 period. The patents used for comparison in this table are those patents for which <http://www.AnyBirthday.com> produced the direct match. (iv) Regressions with technological category controls are reported using the six-category measure of Hall *et al.* (2001). Results using the 36-category.

(b) Table B.3: Inventors per patent, mean differences between samples

8 Short summary

- The paper asks whether innovation is getting harder because the frontier of knowledge becomes more costly to reach.
- It documents three major time trends: inventors enter later, specialize more, and work in larger teams.
- It documents three cross-sectional facts: deeper fields have larger teams and more specialization, but not much higher age at first innovation.
- The data combine a large patent database with inventor-age matches from public records.
- Team size is measured by inventors per patent; specialization is measured by the probability that a solo inventor switches fields; knowledge depth is measured by citation-tree size.
- The model explains why education can be equalized across fields while specialization and teamwork vary across fields.
- The key theoretical mechanism is endogenous sorting across application areas combined with a trade-off between breadth and education cost.
- The model's growth equation shows that the burden of knowledge can eliminate scale effects: more researchers are needed just to sustain growth.
- The paper links the rise of specialization and teams to the decline of the Renaissance Man.

9 Conclusion

9.1 Core conclusion

Jones argues that knowledge accumulation has a double-edged effect. It makes future innovation possible, but it also raises the educational burden that new innovators must bear. The result is a shift in the organization of innovation: individuals become more specialized, teams become larger, and innovators enter active creative work later in life.

9.2 Economic intuition

The intuition is that a single innovator has limited time and learning capacity. As the frontier gets deeper, one person cannot cheaply cover the same breadth of knowledge. The economy responds by dividing knowledge across people. Specialization and teamwork are therefore equilibrium responses to the rising cost of individual breadth.

9.3 Incremental contribution revisited

The paper contributes by connecting micro-level innovator behavior to macro-level growth facts. It does not simply assume that ideas are harder to find. Instead, it shows how the human-capital burden of reaching the frontier can make innovation harder. This mechanism explains why R&D

effort can rise without accelerating productivity growth and why innovative activity becomes more team based.

9.4 Implications for innovation policy and organization

The model gives education policy a central role. If the rate at which society transfers knowledge to new cohorts improves, then the burden of knowledge can be partially offset. Curriculum design, scientific training, graduate education, collaboration tools, and institutions that help people learn faster may therefore have first-order effects on long-run growth.

The paper also implies that firms, universities, and laboratories should expect innovation to become increasingly team based. Effective coordination across specialized experts becomes an essential input into technological progress.

9.5 Limitations

The empirical analysis is based on patents, which do not capture all innovation. The age data are selected because birth dates are available only for a subset of inventors. Citation-tree size is an imperfect proxy for knowledge depth. The specialization measure is cleanest for solo inventors and may not fully capture expertise in team production. Finally, the regressions document patterns that are consistent with the model, but they are not randomized causal estimates.

9.6 Takeaway

The enduring takeaway is that the rise of teams and specialization may reflect a fundamental feature of technological progress. As knowledge accumulates, the frontier becomes harder for any single mind to span. Sustained growth may therefore require ever more researchers, better knowledge transmission, and better institutions for combining specialized expertise.