

Notes: Radner (Econometrica, 1993)

The Organization of Decentralized Information Processing

Contents

1	Big picture	3
2	Incremental contribution of the paper	3
3	Model primitives and measurement	4
3.1	Processors and items	4
3.2	Networks and programs	4
3.3	Hierarchy and ranked trees	4
3.4	Efficiency variables	4
3.5	One-shot versus systolic mode	5
4	Models and theoretical specifications	5
4.1	One-shot minimum-delay problem	5
4.2	Attainability by a tree	5
4.3	Processor-delay cost function	5
4.4	Balanced hierarchies	6
4.5	Preprocessing/tree networks	6
4.6	Asymptotic cost of PPT networks	6
4.7	Systolic-mode lower bounds	6
4.8	Applications: linear rules, project selection, pattern matching	7
5	Theoretical design and identification-style logic	7
5.1	Design object: computational feasibility	7
5.2	Why the one-shot theorem is credible	7
5.3	Why hierarchies are sufficient but not always literal organization charts	8
5.4	Why systolic mode is harder	8
5.5	What is held fixed and what is abstracted from	8
5.6	How to read the efficiency results	8
6	Tables and figures	9
6.1	Figure 1: A tree; Figure 2: The same tree organized in levels in two different ways	9
6.2	Figure 3: A hierarchical network with 15 processors, 40 items, and delay 11	10
6.3	Table I: Minimum Delay (C) for Adding 40 Items; Figure 4: $N = 40$, Min C vs. P	10
6.4	Figure 5: $N = 10,000$, Min C vs. P	11
6.5	Figures 6a–6d: Balanced hierarchy before and after staged reduction	12
6.6	Table II: Relative Inefficiency of Preprocessing/Tree Networks	13

6.7	Figure 7: Efficient systolic-mode networks and the lower bound	13
6.8	Figure 8: Some good balanced trees	14
7	Short summary	14
8	Conclusion	15
8.1	Core conclusion	15
8.2	Economic intuition	15
8.3	Incremental contribution revisited	15
8.4	Policy and management implications	15
8.5	Limitations	15
8.6	Takeaway	16

1 Big picture

- **Question.** How should a large organization process many pieces of information when each manager or processor has limited capacity?
- **Core perspective.** The paper models the managerial side of a firm as a decentralized information-processing machine. Signals enter the organization, processors transform and aggregate them, and a final decision is produced after some delay.
- **Main trade-off.** The organization can reduce delay by using more processors and more parallel processing, but processors are scarce. Efficient organization is therefore a frontier between the number of processors P and the delay C .
- **Key insight on hierarchy.** Hierarchies are often viewed as centralized authority structures, but Radner shows that they can also be efficient devices for decentralizing information processing.
- **One-shot problem.** For a single cohort of N information items, the paper characterizes the minimum delay achievable with P processors and shows that a tree can attain the bound.
- **Repeated problem.** For recurring cohorts arriving every T cycles, the paper studies steady or “systolic” processing, where different cohorts move through the network simultaneously like a pipeline.
- **Organizational production function.** The paper interprets the feasible frontier among N , P , and C as a production function for management: how many items can be processed with how many processors and how much delay.
- **Applications.** The same associative-operations logic applies to linear decision rules, accounting and control, project selection, and pattern recognition.
- **Economic takeaway.** Bounded managerial capacity is not a side detail; it is an input constraint that shapes hierarchy, decentralization, delay, and returns to scale in firms.

2 Incremental contribution of the paper

- **Brings parallel computation into the economics of organization.** Earlier economics of teams and organizations studied information acquisition, communication, and incentives. This paper adds a computational dimension: how long it takes bounded agents to process information.
- **Turns hierarchy into a processing technology.** The paper does not merely describe hierarchy as a chain of command. It shows how tree-like networks aggregate information and where they lie on an efficiency frontier.
- **Measures efficiency with two scarce resources.** Computer science often focuses on time or delay. Radner gives equal attention to the number of processors, which is economically natural because processors are costly managerial inputs.
- **Derives explicit one-shot frontier.** Theorem 1 gives a sharp formula for the minimum delay for processing N items with P processors and shows attainability by a tree.
- **Connects one-shot and repeated regimes.** The paper shows that the one-shot analysis is not an isolated exercise; it becomes the building block for steady-state repeated processing.
- **Clarifies returns to scale in management.** If delay is treated as output quality, then scaling up the number of items can exhibit decreasing returns because constant delay cannot be maintained as N grows unless processors grow sufficiently fast.
- **Prefigures modern bounded rationality and organizational computation.** The paper offers a concrete bounded-rationality technology: agents are rational in design, but individually limited in processing capacity.
- **Provides a bridge between formal organizations and algorithms.** Concepts such as balanced trees, preprocessing, systolic networks, and associative operations are translated into organizational language.

3 Model primitives and measurement

3.1 Processors and items

The organization must process N items of information. A processor is an individual information-processing unit with:

- an in-box;
- a register;
- a clock;
- a finite processing capacity.

In one cycle, a processor can take one item from its in-box and combine it with the contents of its register. The canonical operation is addition, but the paper stresses that the analysis applies to any associative operation.

Interpretation: Associativity is the mathematical reason parallelization is possible. If $(a+b)+c = a+(b+c)$, partial sums can be produced in different parts of the organization and then combined later.

3.2 Networks and programs

A network specifies processors and one-way communication links among them. A program specifies:

- which items enter which processors;
- when each processor sends its register contents to another processor;
- which processor outputs the final answer;
- the timing of messages and aggregation.

The network is not assumed to be a hierarchy, but hierarchies and trees are especially important because they are sufficient for efficient one-shot aggregation.

Interpretation: The distinction between a network and a program matters. A physical organization chart alone does not determine performance; the rules governing timing, routing, and aggregation also matter.

3.3 Hierarchy and ranked trees

A hierarchy is a ranked tree. Processors are assigned levels. The root is at the highest level. Each non-root processor reports to one immediate superior. Data enter at the bottom and partial aggregates flow upward.

Interpretation: In Radner's model, hierarchy is not primarily about authority. It is a computational architecture that transforms many local observations into one aggregate decision.

3.4 Efficiency variables

The central objects are:

- N : number of items in a cohort;
- P : number of processors;
- C : delay, measured as the number of cycles between input and final output;
- T : time between arrivals of cohorts in repeated processing.

A pair (P, C) is efficient for N if one cannot reduce P without increasing C , and cannot reduce C without increasing P .

Interpretation: This is a Pareto frontier in organizational design. Efficient organizations are not necessarily low-delay or low-processor organizations; they are designs where no free improvement is possible along the P – C trade-off.

3.5 One-shot versus systolic mode

- **One-shot mode:** one cohort of N items is processed once.
- **Systolic mode:** a new cohort of N items arrives every T cycles, so several cohorts may be in the network at once.

Interpretation: The one-shot problem is easier and gives sharp results. The repeated problem is more relevant for ongoing firms, where information keeps arriving and decisions must be continuously updated.

4 Models and theoretical specifications

4.1 One-shot minimum-delay problem

Given N items and P processors, the delay-minimization problem asks how quickly the organization can aggregate the N items.

For $N \geq P \geq 1$, Theorem 1 gives the minimum delay:

$$\hat{C}(N, P) = \left\lfloor \frac{N}{P} \right\rfloor + \lceil \log_2 (P + N \bmod P) \rceil. \quad (1)$$

Interpretation: The first term is the serial work each processor must do at the bottom level. The second term is the time needed to aggregate the partial results through a binary-like tree. Adding processors reduces the serial term, but the aggregation term remains.

4.2 Attainability by a tree

Theorem 1 also states that the minimum delay is attainable by a tree. This is a strong result because it says general networks do not outperform tree structures for the one-shot aggregation problem.

Interpretation: Even though the model permits arbitrary directed networks, a hierarchical tree can be efficient. This formalizes the idea that hierarchy can be an optimal decentralized processing architecture.

4.3 Processor-delay cost function

The paper often discusses minimizing a positive linear function of delay and processors:

$$L = \gamma C + \phi P, \quad (2)$$

where $\gamma > 0$ is the cost of delay and $\phi > 0$ is the cost of a processor.

Interpretation: Different environments have different relative prices. If delay is very costly, the organization should use more processors. If managerial processors are expensive, the organization tolerates more delay.

4.4 Balanced hierarchies

A balanced hierarchy has levels $r = 1, \dots, R$, with P_r processors at level r . If the fan-in at level r is k_r , then approximately

$$P_r \approx \frac{P_{r-1}}{k_r}. \quad (3)$$

At the bottom level, each first-level processor receives roughly

$$k_1 \approx \frac{N}{P_1} \quad (4)$$

original items.

Interpretation: The fan-in sequence determines how much serial work occurs at each level. Balanced hierarchies are easy to describe, but one-shot optimal hierarchies can become visually unbalanced after reduction.

4.5 Preprocessing/tree networks

A preprocessing/tree (PPT) network has a first level of preprocessors and then a balanced tree with common fan-in K above the first level. If k_1 is the amount of preprocessing by first-level processors, then roughly:

$$N \approx P_1 k_1. \quad (5)$$

The upper-level tree aggregates the partial results using fan-in K .

Interpretation: PPT networks split the organization into a wide preprocessing stage and a narrower aggregation stage. They are simple and almost efficient in large systems.

4.6 Asymptotic cost of PPT networks

For relatively optimal PPTs, the paper reports the asymptotic relationships:

$$C \sim (\alpha N)^{1/2}, \quad P \sim \left(\frac{N}{\alpha}\right)^{1/2}, \quad (6)$$

$$k_1 \sim (\alpha N)^{1/2}, \quad K \sim \frac{(4\alpha N)^{1/4}}{\ln N}, \quad (7)$$

where $\alpha = \phi/\gamma$.

Interpretation: With fixed relative prices, both delay and the number of processors grow like the square root of the number of items. This is an organizational production function: larger decision problems require more processors and more delay.

4.7 Systolic-mode lower bounds

In systolic mode, a cohort of N items arrives every T cycles. The paper gives lower bounds of the form:

$$P \geq \frac{N}{T}, \quad (8)$$

$$C \geq 1 + \log_2 N. \quad (9)$$

Interpretation: The first bound is a throughput bound: if N items arrive every T cycles, the organization must have enough capacity to process at least N/T items per cycle. The second bound is an aggregation-depth bound: even with massive parallelization, combining many pieces of information takes time.

4.8 Applications: linear rules, project selection, pattern matching

Radner sketches three applications.

- **Linear decision rules.** Accounting and control problems involve multiplying data by weights and adding them.
- **Project selection.** Selecting the best project is a minimum or maximum operation, which is associative under pairwise comparisons.
- **Pattern matching.** Classifying an environmental vector by nearest reference vector can be decomposed into many distance calculations, sums, and a final minimum.

Interpretation: The same organizational logic applies whenever the task can be decomposed into local operations and recombined through an associative operation.

5 Theoretical design and identification-style logic

This is a theoretical paper, not an empirical identification paper. The relevant design question is therefore not how a causal effect is identified from data, but how the model isolates the processing constraints that determine organizational efficiency.

5.1 Design object: computational feasibility

The paper specifies a minimal technology for information processing:

- processors are identical;
- each processor can process one item per cycle;
- links transmit partial results upward or across the network;
- delay is observable and costly;
- processor count is observable and costly.

Interpretation: The stripped-down technology makes the efficiency frontier transparent. Once this frontier is known, richer organizational features can be interpreted as adding constraints or costs.

5.2 Why the one-shot theorem is credible

Theorem 1 has two sides.

- **Lower bound:** no network can complete the task faster than the formula because each processor can process only one item per cycle and aggregation requires sufficient combining depth.
- **Construction:** a tree can attain the bound, proving the lower bound is tight.

Interpretation: The proof strategy mirrors identification in theory: a lower bound shows what is impossible, while a constructive algorithm shows what is attainable.

5.3 Why hierarchies are sufficient but not always literal organization charts

The efficient trees can be visually odd. In reduced balanced hierarchies, high-level processors may have immediate subordinates at many different levels.

Interpretation: The mathematical tree is an information-processing hierarchy, not necessarily a real authority hierarchy. The result says that tree-like aggregation is efficient, not that real firms should literally draw organization charts that look like the reduced tree.

5.4 Why systolic mode is harder

In one-shot mode, only one cohort moves through the network. In systolic mode, many cohorts may be in process simultaneously. This creates pipeline constraints:

- lower levels may work on a new cohort while upper levels finish an old one;
- idle time matters;
- delay and throughput must be optimized together;
- reusing processors across cohorts can be more efficient than replicating one-shot trees.

Interpretation: The repeated-regime problem resembles a production line. Efficient design is not only about finishing one job quickly, but about maintaining flow across many jobs.

5.5 What is held fixed and what is abstracted from

The paper deliberately abstracts from:

- communication costs;
- heterogeneous processor abilities;
- endogenous data acquisition;
- incentive problems and private information;
- errors, attention failures, and strategic behavior;
- multi-dimensional authority or coordination conflicts.

Interpretation: These abstractions are not weaknesses for the narrow question. They allow the paper to isolate pure processing capacity. Later organizational economics can add communication, incentives, and errors on top of the computational baseline.

5.6 How to read the efficiency results

An efficiency frontier in this paper should be read as:

$$\text{tasks processed } (N) \longleftrightarrow \text{processors } (P) \longleftrightarrow \text{delay } (C). \quad (10)$$

Interpretation: A firm with more managers can process more information or respond faster, but the frontier says there is no free lunch. More scale requires more managerial capacity, more delay, or both.

6 Tables and figures

6.1 Figure 1: A tree; Figure 2: The same tree organized in levels in two different ways




FIGURE 1.—A tree.

(a) **Figure 1: A tree.** The basic object is a rooted tree. The root is at the top, and not every pair of objects is comparable. This figure separates tree structure from ranked hierarchy.




FIGURE 2.—The same tree organized in levels in two different ways.

(b) **Figure 2: The same tree organized in levels in two different ways.** The same underlying tree can be assigned different rankings or levels, showing that hierarchy requires both a tree and a ranking.

Read:

- Figure 1 introduces the primitive graph object used throughout the paper.
- Figure 2 shows that hierarchy is not just graph structure; it also requires a level assignment.
- The figures clarify why a processor can be neither above nor below another processor even inside the same organization.

Detailed interpretation: The difference between a tree and a hierarchy is important because the paper later treats levels as processing stages. If the same tree can be ranked in different ways, then a fixed reporting relationship does not uniquely determine processing delay. The organization also needs rules assigning work to levels and scheduling communication.

Economic message: A firm chart is not enough. What matters is how the chart is programmed as an information-processing device.

6.2 Figure 3: A hierarchical network with 15 processors, 40 items, and delay 11




FIGURE 3.—A hierarchical network with 15 processors, 40 items, and delay 11.

Figure 3: A hierarchical network with 15 processors, 40 items, and delay 11. The 40 items enter at the bottom through first-level processors. Partial sums are sent upward until the root outputs the grand total.

Read:

- The network uses 15 processors.
- The 40 items are divided among the lowest-level processors.
- Processors at the same level work in parallel, but each processor processes serially.
- The example takes 11 cycles, even though the later efficiency formula shows that the same N can be processed in fewer cycles.

Detailed interpretation: This figure is the organizational analogue of a parallel addition circuit. The bottom level performs wide parallel preprocessing, while upper levels aggregate partial sums. The example is intuitive but not efficient: it uses a balanced-looking hierarchy that leaves room for reduction.

Economic message: A natural-looking hierarchy may waste managerial resources. Efficient information processing can require restructuring the hierarchy or changing the program.

6.3 Table I: Minimum Delay (C) for Adding 40 Items; Figure 4: $N = 40$, Min C vs. P

TABLE I
MINIMUM DELAY (C) FOR ADDING 40 ITEMS^a

P	1	2	3	4	5	6	7	8	9	10	11	12	13	...	20	...	40
Min C	40	21	15	12	11	10	9	8	8	8	8	7	7	...	7	...	7
	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*	*

^a P = number of processors, * denotes an efficient network.

(a) **Table I: Minimum Delay (C) for Adding 40 Items.** The table reports the minimum delay for $N = 40$ as the number of processors P changes. Asterisks denote efficient networks.




FIGURE 4.— $N = 40$. Min C vs. P .

(b) **Figure 4: $N = 40$. Min C vs. P .** The plot visualizes diminishing delay reductions as processors are added.

Read:

- With one processor, delay is 40 cycles.
- With two processors, delay falls to 21 cycles.
- The minimum delay is 7 cycles and is attainable with 12 processors.
- More than 12 processors do not reduce delay below 7 cycles, so extra processors are inefficient for this N .
- Some processor counts are dominated: for example, 9–11 processors give delay 8, but 8 processors can also achieve delay 8.

Detailed interpretation: Table I is the cleanest numerical demonstration of the P – C frontier. The delay curve falls sharply at first because additional processors reduce serial work at the bottom level. Eventually the aggregation depth becomes the bottleneck, so adding processors stops improving delay.

Economic message: The efficient organization is not the one with the most managers. Once delay cannot fall further, additional managerial processors are redundant.

6.4 Figure 5: $N = 10,000$, Min C vs. P




FIGURE 5.— $N = 10,000$. Min C vs. P .

Figure 5: $N = 10,000$. Min C vs. P . The figure shows the delay frontier on a larger scale. Inefficiency from processor-count discreteness becomes relevant only when P is quite large.

Read:

- For large N , delay falls rapidly as processors are added at low P .
- The curve flattens as P becomes large.
- For $N = 10,000$, Radner notes that inefficiency is not a major issue until the number of processors exceeds roughly 2,000.

Detailed interpretation: The large- N plot shows that the one-shot frontier behaves smoothly for most practical processor counts. Discrete jumps matter, but the central economic lesson remains the serial-parallel trade-off.

Economic message: Large organizations can gain a lot from decentralizing processing at first, but the returns to additional decentralization diminish.

6.5 Figures 6a–6d: Balanced hierarchy before and after staged reduction




FIGURE 6a.—Balanced hierarchy before reduction.




FIGURE 6b.—Balanced hierarchy after 1 stage of reduction.




FIGURE 6c.—Balanced hierarchy after 2 stages of reduction.




FIGURE 6d.—Balanced hierarchy after 3 stages of reduction (final).

Figures 6a–6d: Balanced hierarchy before reduction; after 1 stage; after 2 stages; after 3 stages. The figure shows how a balanced hierarchy can be reduced by eliminating processors and reassigning their predecessors upward.

Read:

- Figure 6a starts from a balanced hierarchy similar to Figure 3.
- Figure 6b eliminates one processor from each first-level cadre.
- Figures 6c and 6d continue the reduction at higher levels.
- The number of processors falls from 15 to 8.
- The delay falls from 11 to 8 cycles.

Detailed interpretation: The result is counterintuitive: removing processors can reduce both the number of processors and delay. The reason is that a balanced hierarchy may allocate too much work to intermediate processors and generate unnecessary levels. Reduction reallocates predecessors to make the tree closer to the efficient one-shot design.

Economic message: Organizational simplification can improve speed and reduce managerial headcount simultaneously. This is a formal argument for why layers of hierarchy may be inefficient even if the organization appears balanced.

6.6 Table II: Relative Inefficiency of Preprocessing/Tree Networks

TABLE II RELATIVE INEFFICIENCY OF PREPROCESSING / TREE NETWORKS				
N	10^4	10^6	10^8	10^{10}
P	10^2	10^3	10^4	10^5
$\hat{C}$	107	1010	10014	100017
$\hat{C}_{PPT}$	131	1095	10294	100922
K	10	27	79	242
$\frac{\hat{C}_{PPT}}{\hat{C}}$	1.224	1.084	1.028	1.009

Table II: Relative Inefficiency of Preprocessing/Tree Networks. The table compares optimal tree delay with preprocessing/tree delay for large N and selected P .

Read:

- The ratio $\hat{C}_{PPT}/\hat{C}$ is 1.224 for $N = 10^4$.
- The ratio falls to 1.084 for $N = 10^6$.
- It falls further to 1.028 for $N = 10^8$ and 1.009 for $N = 10^{10}$.
- PPT networks become relatively close to optimal as N grows.

Detailed interpretation: Table II is important because it shows that simple, regular networks can be nearly efficient in large-scale problems. A fully optimal tree may be complicated and unbalanced, but a PPT network has a simple organizational interpretation and loses little delay efficiency asymptotically.

Economic message: Real firms may choose simple hierarchical routines rather than exact optimum networks. Table II explains why this can be sensible: simple designs can be close to optimal at scale.

6.7 Figure 7: Efficient systolic-mode networks and the lower bound

Read:

- The vertical axis is $\log C$, the log of delay.
- The horizontal axis is $\log(P - 10^5)$, since the throughput bound requires $P \geq N/T = 10^5$.
- The efficient systolic networks lie close to the lower bound.
- Adding processors above the throughput minimum reduces delay, but the frontier again flattens.

Detailed interpretation: Figure 7 shifts from processing one cohort to processing a continuous stream of cohorts. It shows that even in a dynamic pipeline, there is a frontier between additional processors and shorter delay. The lower bound is not just theoretical decoration; efficient networks lie near it.

Economic message: In ongoing organizations, the question is not only how fast one report can be processed, but how many reports can be processed continuously without backlog. Systolic networks formalize throughput management.




FIGURE 7.—Efficient systolic-mode networks and the lower bound.

Figure 7: Efficient systolic-mode networks and the lower bound. For $N = 10^6$ and $T = 10$, the figure plots efficient repeated-processing networks against a lower bound.




FIGURE 8.—Some good balanced trees.

Figure 8: Some good balanced trees. The figure compares preprocessing/overhead networks, Garey trees, and the lower bound for $N = 10^6$ and $T = 10$.

6.8 Figure 8: Some good balanced trees

Read:

- Solid dots represent preprocessing/overhead networks.
- Plus signs represent Garey trees.
- The solid curve is the lower bound.
- Preprocessing/overhead networks are often close to the lower bound, especially away from the extreme low-delay region.
- Garey trees perform well when delay is close to the minimum.

Detailed interpretation: Figure 8 is a design menu. It shows that different simple families of balanced trees are good in different regions of the frontier. No single organizational architecture is uniformly best for every relative price of processors and delay.

Economic message: The optimal organization depends on what the firm values more: conserving managerial processors or reducing response delay. This is why organizational design should vary across environments.

7 Short summary

- The paper models a firm as an information-processing organization with bounded processors.
- The key scarce inputs are processors and time.
- In one-shot mode, the paper gives a sharp formula for the minimum delay to aggregate N items with P processors.

- Efficient one-shot processing can be attained by a tree.
- Natural balanced hierarchies can be inefficient, but reduction procedures can lower both delay and processor count.
- Repeated or systolic processing introduces throughput constraints because cohorts arrive continuously.
- PPT networks and other balanced tree families can be close to efficient for large-scale repeated processing.
- The paper interprets the feasible set among N , P , and C as an organizational production function.
- The model applies to any associative operation, including addition, maximum/minimum search, and nearest-neighbor classification.
- The main economic message is that managerial hierarchy can be understood as a technology for decentralizing computation under bounded capacity.

8 Conclusion

8.1 Core conclusion

Radner's central conclusion is that the organization of information processing is a core economic problem. Large firms require many managers not simply because authority must be delegated, but because information must be processed by bounded human agents. Efficient organization is therefore governed by a trade-off between the number of processors and decision delay.

8.2 Economic intuition

The intuition is that processing many signals requires both parallel and serial work. Parallel work spreads items across many processors; serial work remains because each processor has limited capacity and partial results must eventually be combined. Hierarchies are useful because they permit local aggregation followed by higher-level aggregation.

8.3 Incremental contribution revisited

The paper's lasting contribution is to make computational capacity an explicit input into the theory of the firm. It imports ideas from parallel algorithms into organizational economics, but translates them into economic variables: managers, delay, throughput, and returns to scale. This gives a precise way to think about information overload and managerial bottlenecks.

8.4 Policy and management implications

The model suggests that organizational redesign can improve performance even without changing incentives or worker skills. Removing redundant layers, changing routing rules, or altering the balance between preprocessing and aggregation can change both delay and staffing requirements.

8.5 Limitations

The model abstracts from communication costs, incentives, heterogeneity, errors, endogenous information acquisition, and strategic behavior. These omissions are important in real firms. However, the abstraction is useful because it isolates one mechanism: the computational cost of management.

8.6 Takeaway

The paper provides a computational foundation for decentralization. Hierarchy is not only a mechanism of control; it is a technology for processing information at scale. Once managerial attention is scarce, organizational form becomes a production technology with its own efficiency frontier.