

Notes: Zaffaroni (JPE, 2025)

Factor Models for Conditional Asset Pricing

Contents

1	Big picture	3
2	Data and measurement (key objects)	3
2.1	Observed return panel	3
2.2	Unobservables and timing	4
2.3	Constructed objects and empirical exercises	4
3	Models and empirical specifications (all equations + interpretation)	5
3.1	Conditional factor model	5
3.2	Local approximation over a rolling window	5
3.3	The key local-PCA identification logic	6
3.4	Local PCA estimator	6
3.5	What fixed T identifies: ex post rather than ex ante objects	6
3.6	Estimating the number of factors	7
3.7	Testing pricing ability	7
3.8	Mean-variance portfolios and local strong factors	8
3.9	Monte Carlo benchmark against IPCA	8
4	Identification and assumptions (what makes the estimates credible)	8
4.1	What fails in the standard approach	8
4.2	What identifies the model in the new approach	9
4.3	What fixed T does not identify	9
4.4	Why the pricing tests are still valid	9
4.5	Relation to the literature	9
5	Tables and figures	10
5.1	Table 1: Pricing Ability — Local PCA versus IPCA	10

5.2	Table 2: Weak Factors — Local PCA versus Constant PCA	10
5.3	Table 3: Pricing Ability — Local PCA versus Constant PCA	11
5.4	Comment on figures	11
6	Short summary	12
7	Conclusion	12
7.1	Core conclusion (what the paper establishes)	12
7.2	Economic intuition	12
7.3	Takeaway	12

1 Big picture

- **Question.** How can we do reliable inference for *conditional* asset-pricing factor models when the model may omit relevant risk factors and when the conditional dynamics of returns may be misspecified?
- **Setting.** The paper studies a large panel of individual asset returns. The cross section is large (N large), but the relevant time window for estimation is intentionally short (T fixed and possibly small), so the method is explicitly local in time.
- **Core challenge.** Standard PCA-based asset-pricing methods usually assume parameter constancy and rely on asymptotics in which both N and T grow. That is not suitable for conditional models with time-varying betas, time-varying risk premia, and even a time-varying number of factors.
- **Main methodological idea.** Estimate factors by **local PCA**: apply PCA to a rolling window of fixed length T , so the factor model is only required to be approximately constant *within* each local window rather than over the full sample.
- **First contribution.** The paper derives the asymptotic theory of local PCA for conditional factor models when $N \rightarrow \infty$ and T is fixed.
- **Second contribution.** It shows how to do valid inference for the main asset-pricing objects—risk premia, the stochastic discount factor (SDF), and mean-variance portfolio returns—despite the fact that factor loadings are not consistently estimated when T is fixed.
- **Third contribution.** It develops formal specification tests for the pricing ability of the *ex ante* risk premia and *ex ante* SDF, even though only noisy *ex post* counterparts are point-identified in finite- T windows.
- **Main quantitative message.** In Monte Carlo exercises, local PCA performs well relative to IPCA, especially when instruments are misspecified or when betas vary over time. In the empirical application to US monthly stock returns, local PCA typically achieves smaller pricing errors than constant PCA and can recover locally strong factors that constant PCA tends to wash out.
- **Why this paper matters.** It gives a general, nonparametric, rolling-window framework for conditional asset pricing that is robust to omitted factors and flexible enough to accommodate local time variation in the factor structure.

2 Data and measurement (key objects)

2.1 Observed return panel

- **Unit of observation.** Asset i in time period s .
- **Main observable.** Excess return x_{is} for asset i at date s .
- **Risk-free asset.** The model allows for a tradable risk-free asset with return r_{s-1}^f .

- **Asymptotic environment.** The paper works with a **large cross section** and a **fixed local time window**. The object of interest is not the full-sample unconditional factor structure but the factor structure over a short interval.
- **Empirical application.** Monthly US stock returns from CRSP, January 1980 to December 2019, with $N = 19,259$ stocks in the full unbalanced panel. The local window is $T = 24$ months.
- **Balanced versus unbalanced panels.** For direct comparison with constant PCA, the balanced panel has $N = 283$ stocks. The unbalanced local-PCA exercise uses between 1,820 and 2,360 assets per window when $T = 24$.

2.2 Unobservables and timing

- **Latent factors.** F_s is the vector of latent risk factors.
- **Risk premia.** $g_{s-1} = \mathbb{E}_{s-1}(F_s)$.
- **Time-varying loadings.** $k_{i,s-1}$ is the asset's exposure to the factors, known at time $s - 1$.
- **Pricing error / intercept.** $a_{i,s-1}$ is a time-varying intercept, interpreted in asset pricing as a pricing error.
- **Idiosyncratic shock.** e_{is} is the idiosyncratic return component.
- **Time-varying factor dimension.** The number of factors, r_{s-1} , is itself allowed to vary over time, though it is assumed to be locally constant within each estimation window.
- **Timing.** All conditional moments are taken with respect to the information set available at time $s - 1$.

2.3 Constructed objects and empirical exercises

- **Estimated local factors.** $\tilde{F}_s$ from rolling-window PCA.
- **Estimated ex post risk premia.** The time average of the local factors over the current window.
- **Estimated ex post SDF.** Constructed from the local factors and their sample covariance matrix.
- **Pricing statistics.** Bias-corrected cross-sectional averages of estimated pricing errors, $N^{-1} \sum_i \tilde{d}_i^2$ and $N^{-1} \sum_i \tilde{h}_i^2$.
- **Monte Carlo comparisons.** Local PCA is compared with IPCA and with constant PCA.
- **Empirical comparison.** Local PCA and constant PCA are compared on in-sample and out-of-sample pricing performance for risk premia and the SDF.

3 Models and empirical specifications (all equations + interpretation)

3.1 Conditional factor model

The paper starts from a conditional linear factor model for excess returns:

$$x_{is} = a_{i,s-1} + k'_{i,s-1} F_s + e_{is}.$$

- **Interpretation.** Excess returns are driven by a time-varying intercept, a time-varying exposure to latent factors, and an idiosyncratic shock.
- **Conditional risk premia.** The factor means are

$$g_{s-1} = \mathbb{E}_{s-1}(F_s).$$

So the model is conditional because both factor exposures and factor premia may change over time.

A central asset-pricing implication of asymptotic no-arbitrage is the square-summability of pricing errors:

$$\sum_{i=1}^{\infty} a_{i,s-1}^2 < \infty.$$

- **Interpretation.** Cross-sectionally, pricing errors are small enough that they do not create asymptotic arbitrage opportunities.
- **Why it matters econometrically.** This condition lets the paper avoid de-meaning returns. That is crucial because de-meaning would induce problematic serial dependence when T is fixed.

3.2 Local approximation over a rolling window

Fix a date t and let the estimation window be

$$\mathcal{T}_t = \{t - T + 1, \dots, t\}.$$

Within this window, the paper imposes only *local* constancy:

$$r_s = r_{t-1} \quad \text{for all } s \in \mathcal{T}_t,$$

and a smoothness condition on the loadings,

$$\sup_{s \in \mathcal{T}_t} \|\Lambda_s - \Lambda_{t-1}\|^2 = o(N^{1/2}).$$

Writing the model in matrix form over the window:

$$X = a + F\Lambda' + e.$$

- **Interpretation.** The model treats the factor structure as approximately constant only over a short local window rather than over the full sample.
- **Economic meaning.** This is how the paper permits model-free time variation in betas, risk premia, idiosyncratic risk, and even the number of factors.

3.3 The key local-PCA identification logic

The local covariance matrix of returns satisfies

$$\frac{XX'}{NT} \rightarrow_p \frac{1}{T} F \Sigma_\Lambda F' + \frac{1}{T} \Sigma_e.$$

Here Σ_Λ is the cross-sectional second moment of factor loadings and Σ_e is the cross-section average covariance matrix of the idiosyncratic shocks over the local time window.

The crucial assumption is that, after cross-sectional averaging,

$$\Sigma_e = j_t^2 I_T.$$

- **Interpretation.** Once the averaged idiosyncratic covariance matrix is proportional to the identity, its eigenvectors do not contaminate the factor space. The top eigenvectors of XX'/NT then recover the local factor space.
- **Why large N matters more than large T .** The cross section is what disciplines the factor estimation. Accurate factor extraction needs a large number of assets, not a long time series.

3.4 Local PCA estimator

The estimator solves the least-squares problem

$$(\tilde{F}, \tilde{\Lambda}) = \underset{F, \Lambda}{\operatorname{argmin}} \frac{1}{NT} \sum_{s \in \mathcal{T}_t} \sum_{i=1}^N (x_{is} - F'_s k_i)^2.$$

Operationally, $\tilde{F}$ is obtained from the r_{t-1} largest eigenvectors of XX'/NT , multiplied by $\sqrt{T}$, and the loading estimator is

$$\tilde{\Lambda} = X' \tilde{F} / T.$$

Theorem 1 gives the core asymptotic result:

$$\sqrt{N} (\tilde{F}_s - \tilde{H}'_{t-1} F_s) \Rightarrow \mathcal{N}(0, \cdot),$$

where $\tilde{H}_{t-1}$ is the rotation matrix.

- **Interpretation.** The local factors are estimated consistently up to rotation, and their estimation error is asymptotically normal when $N \rightarrow \infty$.
- **Important limitation.** The loadings themselves are not consistently estimated when T is fixed. But their *cross-sectional moments* can still be estimated after bias correction.

3.5 What fixed T identifies: ex post rather than ex ante objects

With a fixed window length, the paper shows that one can identify only an **ex post** version of factor premia,

$$\tilde{F} = \frac{1}{T} \sum_{s \in \mathcal{T}_t} \tilde{F}_s,$$

not the true ex ante conditional mean $\mathbb{E}_{s-1}(F_s)$.

Likewise, the ex post SDF is constructed as

$$\tilde{M}_{s-1,s} = \frac{1}{r_{s-1}^f} - \frac{1}{r_{s-1}^f} \tilde{F}' \tilde{\Omega}^{-1} (\tilde{F}_s - \tilde{F}),$$

where

$$\tilde{\Omega} = \frac{1}{T} \sum_{s \in \mathcal{T}_t} (\tilde{F}_s - \tilde{F})(\tilde{F}_s - \tilde{F})'.$$

- **Interpretation.** The paper is honest that finite- T windows do not let us identify the ex ante factor moments cleanly.
- **But why this is still useful.** These ex post quantities remain asymptotically valid for pricing, portfolio construction, and formal model comparison once the fixed- T noise is properly accounted for.

3.6 Estimating the number of factors

Let $V(k)$ be the residual variance from a k -factor local PCA. The paper proposes the fixed- T criterion

$$\text{PCT}_t(k) = \frac{T}{T-k} V(k) + kg(N).$$

The estimated number of factors is

$$\tilde{r} = \underset{0 \leq k \leq T}{\text{argmin}} \text{PCT}_t(k).$$

Theorem 2 shows consistency if

$$g(N) \rightarrow 0, \quad \sqrt{N} g(N) \rightarrow \infty.$$

- **Interpretation.** The penalty shrinks slowly enough to let the signal dominate, but fast enough to rule out spurious small eigenvalues.
- **Why this matters.** Existing factor-number criteria typically require both N and T to diverge. This one is designed for the paper's local fixed- T setup.

3.7 Testing pricing ability

The ex ante pricing errors associated with risk premia and the SDF are defined as

$$d_{i,s-1} = \mathbb{E}_{s-1}(x_{is}) - k'_{i,s-1} \mathbb{E}_{s-1}(F_s), \quad h_{i,s-1} = \mathbb{E}_{s-1}(m_{s-1,s} x_{is}).$$

The estimated counterparts $\tilde{d}_i$ and $\tilde{h}_i$ are contaminated by fixed- T noise because they are built from ex post factor moments. The key insight of Theorem 3 is that the finite- T bias in the cross-sectional averages

$$N^{-1} \sum_{i=1}^N \tilde{d}_i^2, \quad N^{-1} \sum_{i=1}^N \tilde{h}_i^2$$

can be consistently estimated and removed. The resulting bias-adjusted test statistics, $S_{d,t-1}$ and $S_{h,t-1}$, are asymptotically standard normal under correct model specification and diverge under misspecification.

- **Interpretation.** The paper does not confuse fixed- T noise with genuine pricing failure.
- **Economic meaning.** This is what lets the author evaluate whether the *ex ante* risk premia and SDF truly price the cross section.

3.8 Mean-variance portfolios and local strong factors

Even though portfolio weights themselves remain noisy because loadings are noisy, mean-variance *portfolio returns* can still be estimated accurately by exploiting cross-sectional averaging. This also yields valid inference for the Sharpe ratio and value at risk.

The paper additionally defines a **locally strong factor** as one whose squared risk exposures satisfy

$$N^{-1} \sum_{i=1}^N (\lambda_{i,s-1}^{\text{lst}})^2 = \begin{cases} O(1), & s \in \mathcal{T}_t, \\ O(N^{-1}), & s \notin \mathcal{T}_t. \end{cases}$$

- **Interpretation.** A factor can be important over a short historical episode and essentially irrelevant outside that window.
- **Why local PCA matters.** Constant PCA can average such factors away and mistakenly classify them as weak or useless.

3.9 Monte Carlo benchmark against IPCA

To compare local PCA with IPCA, the paper simulates

$$x_{is} = k_{i,s-1} f_s + e_{is}, \quad k_{i,s-1} = z'_{i,s-1} \Gamma + n_i,$$

with latent factors f_s , observed instruments $z_{i,s-1}$, and an unobserved component n_i .

- **Interpretation.** This is a natural benchmark because IPCA is explicitly designed for a characteristics-based conditional factor model.
- **What the paper varies.** It considers the correct instrument set, a misspecified instrument set, and random-walk instruments that generate strong time variation in exposures.

4 Identification and assumptions (what makes the estimates credible)

4.1 What fails in the standard approach

- **Constant PCA is too rigid.** If risk premia, betas, idiosyncratic risk, or the factor dimension vary over time, full-sample constant PCA imposes the wrong restriction.

- **Large- T asymptotics are unavailable.** The classical PCA inferential theory does not apply when the estimation window is fixed.
- **Instrument-based conditional models can be misspecified.** Methods that force time variation to depend on preselected observables may fail when the wrong instruments or the wrong functional form are used.
- **Demeaning is not innocuous.** In the local fixed- T environment, de-meaning induces temporal dependence in the error term and destroys the key covariance structure that local PCA exploits.

4.2 What identifies the model in the new approach

- **Large cross-sectional information.** Factor extraction is driven by $N \rightarrow \infty$.
- **Local smoothness.** Within a short window, the number of factors and the loadings are approximately constant.
- **Cross-sectionally averaged idiosyncratic covariance is scalar.** This allows the eigenvectors of the return covariance matrix to identify the factor space.
- **Distinct eigenvalues.** The factor-related eigenvalues must be separated from the noise eigenvalues.
- **No-arbitrage.** The square-summability of pricing errors eliminates the need to estimate intercepts and makes local PCA statistically feasible.

4.3 What fixed T does not identify

- The *ex ante* risk premia and *ex ante* SDF are not point-identified when T is fixed.
- The factor loadings themselves are not consistently estimated.
- What *is* identified are the factors up to rotation, the moments of the loadings after bias correction, and the ex post objects needed for pricing and portfolio analysis.

4.4 Why the pricing tests are still valid

- The paper estimates the finite- T bias inside $N^{-1} \sum_i \tilde{d}_i^2$ and $N^{-1} \sum_i \tilde{h}_i^2$ rather than ignoring it.
- Once that bias is removed, the remaining statistic isolates genuine pricing performance.
- The resulting tests are invariant to asset repackaging under the null, which is a desirable asset-pricing property.

4.5 Relation to the literature

- The paper is closest in spirit to the PCA-based asset-pricing literature, especially Pelger, Lettau–Pelger, Kelly–Pruitt–Su, Ma–Gagliardini, and Giglio–Xiu.
- Relative to constant PCA, the contribution is to permit model-free local time variation and to derive valid fixed- T inference.

- Relative to IPCA-style methods, the contribution is robustness to misspecified instruments and omitted-factor problems.
- Relative to the conditional beta-pricing literature, the contribution is to test the pricing ability of risk premia and the SDF in a large- N , locally time-varying environment.

5 Tables and figures

5.1 Table 1: Pricing Ability — Local PCA versus IPCA

Read:

- The table reports the ratio of the IPCA pricing criterion to the local-PCA pricing criterion. Values above one mean local PCA prices better.
- Under the **correct stationary-instrument model**, local PCA already dominates in sample: for example, the ratio is 5.851 when $(N, T) = (15, 5)$ and 5.304 when $(N, T) = (1,000, 5)$. Out of sample the comparison is closer, but local PCA is still usually better.
- Under the **misspecified instrument model**, IPCA deteriorates sharply. Out-of-sample ratios range roughly from 1.967 to 3.401 for the smaller designs and stay above one throughout. This is exactly where local PCA's robustness to instrument misspecification matters most.
- Under **random-walk instruments**, IPCA also deteriorates because it struggles with strong time variation in risk exposures. Here too local PCA performs better, especially when N is large.
- Main message: local PCA is competitive when the instrumented model is correct and clearly better when the instrument-based specification is wrong or too rigid.

TABLE 1
PRICING ABILITY: LOCAL PCA VERSUS IPCA

(N, T)	IPCA/Local PCA					
	(15, 5)	(15, 25)	(100, 5)	(100, 25)	(1,000, 5)	(1,000, 25)
A. Correct Model: Stationary Instruments						
In sample	5.851	2.036	5.213	1.479	5.304	2.150
Out of sample	1.137	1.009	1.454	1.487	1.223	2.130
B. Misspecified Model: Stationary Instruments						
In sample	11.668	3.891	6.997	2.764	10.023	3.204
Out of sample	2.145	3.401	1.967	2.002	2.574	3.012
C. Correct Model: Random Walk Instruments						
In sample	9.980	4.568	9.025	3.356	11.655	5.924
Out of sample	1.681	3.494	1.942	2.849	2.077	4.436

NOTE.—The table presents ratios of the pricing ability criterion $N^{-1} \sum_{t=1}^T \hat{\sigma}_t^2$ corresponding to the IPCA (numerator) and local PCA (denominator), averaged across 1,000 Monte Carlo iterations and evaluated in sample and out of sample, for various combinations of (N, T) . The data generating process is (15), where the vector $\mathbf{z}_{t-1}$ is a 10-sized stationary vector autoregression (experiment A) and a multivariate random walk (experiment C), respectively. Experiment B differs from experiment A because the IPCA-estimated model is misspecified (one instrument instead of the 10 required instruments).

TABLE 2
WEAK FACTORS: LOCAL PCA VERSUS CONSTANT PCA ($T = 24, \tau_{t-1} = 2$)

Factor	$\text{cor}(\tilde{\mathbf{F}}_t, \tilde{\mathbf{H}}_{t-1} \mathbf{F}_t)$	$\text{cor}(\tilde{\mathbf{F}}_t^0, \tilde{\mathbf{H}}^0 \mathbf{F}_t)$	$\left(\frac{\tilde{\mathbf{H}}_{t-1}^0 \tilde{\mathbf{F}}_t}{\tilde{\mathbf{F}}_t} - 1\right)\%$	$\left(\frac{\tilde{\mathbf{H}}_{t-1}^0 \tilde{\mathbf{F}}_t}{\tilde{\mathbf{F}}_t} - 1\right)\%$	$\left(\frac{\tilde{\mathbf{H}}_{t-1}^0 \tilde{\mathbf{F}}_t}{\tilde{\mathbf{H}}_{t-1}^0 \tilde{\mathbf{F}}_t} - 1\right)\%$
	(1)	(2)	(3)	(4)	(5)
1	.99	.02	.03	−86.65	3.54
2	.99	.11	−.51	−212.79	26.82

NOTE.—The table reports the correlation between the local PCA risk factor ($\tilde{\mathbf{F}}_t$) and the true (rotated) risk factor and between the constant PCA risk factor ($\tilde{\mathbf{F}}_t^0$) and the true (rotated) risk factor (averaged across T and $T + T^0$ periods, respectively), the bias (in percent) of the local PCA risk premia and the constant PCA risk premia, and the ratio between local PCA and constant PCA cross-averaged squared (rotated) risk exposures (Monte Carlo averages; $M = 1,000$ replications). Rows 1 and 2 correspond to the first and second risk factors, respectively.

5.2 Table 2: Weak Factors — Local PCA versus Constant PCA

Read:

- This table studies a factor that is strong only locally. For both factors, the correlation between the local-PCA estimate and the true rotated factor is 0.99.

TABLE 3 PRICING ABILITY: LOCAL PCA VERSUS CONSTANT PCA											
Constant PCA/Local PCA	$\tilde{r}$	Risk Factor									
		1	2	3	4	5	6	7	8	9	10
		A. SDF ($T = 24$)									
Out of sample:											
Balanced	5.262	27.791	9.503	8.116	4.875	3.106	2.302	1.871	1.323	.915	.712
Unbalanced	1.096	5.416	3.003	1.776	3.329	4.38	1.745	1.468	.581	.864	.721
In sample:											
Balanced	1.870	12.83	4.475	3.213	2.198	1.683	1.323	1.089	.889	.740	.613
Unbalanced	.849	3.510	2.079	1.333	1.082	.804	.633	.513	.429	.365	.298
B. Risk Premia ($T = 24$)											
Out of sample:											
Balanced	.845	10.470	6.567	4.784	3.674	2.927	2.369	1.955	1.631	1.375	1.169
Unbalanced	1.492	5.534	3.407	2.461	1.880	1.493	1.205	.992	.827	.697	.592
In sample:											
Balanced	4.634	23.767	8.151	6.337	5.0468	3.550	3.720	2.465	2.660	2.420	1.738
Unbalanced	2.908	161.306	10.825	5.052	3.246	3.148	3.210	2.094	1.455	1.280	1.330

NOTE.—Each entry represents the ratio between the pricing statistics for the constant PCA over the local PCA for both the balanced and unbalanced panel case: ratios (in absolute value) of the time averages (across 480 $-T$ windows) of the bias-adjusted SDF pricing statistics $N^{-1}\sum_{t=1}^T \hat{A}_t^c / \hat{A}_t^l$ (panel A) and of the bias-adjusted risk premia pricing statistics $N^{-1}\sum_{t=1}^T \hat{B}_t^c / \hat{B}_t^l$ (panel B), both described in theorem 3. Each panel reports both in-sample and out-of-sample results corresponding to the estimated r , described in app. sec. OA15.1, and fixed risk factors ranging from 1 to 10 (monthly data from January 1980 to December 2019; $T = 24$, $N = 19,259$, sourced from CRSP).

- Constant PCA performs very poorly in the same design: the corresponding correlations are only 0.02 for factor 1 and 0.11 for factor 2.
- The local-PCA bias in risk premia is tiny: about 0.03% for factor 1 and -0.51% for factor 2. Constant PCA instead has huge negative biases of about -86.65% and -212.79% .
- Column 5 shows that the second factor appears about 26.82 times weaker under constant PCA than under local PCA.
- Main message: if a factor is only temporarily important, constant PCA can make it look weak or useless, whereas local PCA recovers it almost perfectly.

5.3 Table 3: Pricing Ability — Local PCA versus Constant PCA

Read:

- Each entry is the ratio of the constant-PCA pricing statistic to the local-PCA pricing statistic. Ratios above one indicate smaller pricing errors under local PCA.
- The gains are especially strong in the **SDF pricing** exercise when the factor dimension is kept small. For one factor, the out-of-sample balanced ratio is 27.791 and the in-sample balanced ratio is 12.83.
- The **risk-premia pricing** results also favor local PCA strongly for parsimonious specifications. With one factor, the out-of-sample balanced ratio is 10.470 and the in-sample balanced ratio is 23.767; in the unbalanced in-sample case it is as high as 161.306.
- When the estimated number of factors $\tilde{r}$ is used, the picture is more mixed in some cells, but local PCA still often improves materially, especially for SDF pricing and for the richer unbalanced panels.
- Main message: allowing factor structure to vary locally carries real economic value. The benefits are largest when constant PCA is forced to summarize a time-varying world with a fixed full-sample factor structure.

5.4 Comment on figures

- The main paper is primarily **table-driven**. The central quantitative evidence is carried by Tables 1–3 rather than by headline figures.
- The paper’s logic therefore runs from theory and asymptotics to Monte Carlo tables and then to a direct empirical horse race between local and constant PCA.

6 Short summary

- **Conclusion.** The paper builds a rolling-window PCA methodology for conditional asset pricing with time-varying factor structure and fixed local time windows.
- **Identification insight.** A large cross section identifies latent factors locally, while no-arbitrage and a scalar cross-section average idiosyncratic covariance make the local eigenstructure informative. Fixed- T noise is then handled analytically rather than ignored.
- **Empirical lesson.** Relative to constant PCA and IPCA, local PCA is more robust to omitted factors, instrument misspecification, and locally strong factors, and it typically produces smaller pricing errors in US stock-return data.

7 Conclusion

7.1 Core conclusion (what the paper establishes)

- The paper establishes that conditional asset-pricing factor models can be estimated and tested with a **local PCA** procedure that is asymptotically valid when the cross section is large and the time window is fixed.
- It shows that accurate factor extraction does not require a long time series; what it requires is a rich cross section.
- It also shows that one can do valid inference on risk premia, the SDF, and mean-variance portfolio returns even though factor loadings are not consistently estimated in fixed- T windows.
- Finally, it provides formal pricing tests that separate genuine model misspecification from finite-window estimation noise.

7.2 Economic intuition

- Constant PCA averages across time and therefore tends to wash out local variation in the factor structure.
- Local PCA instead asks a narrower question: *what factor structure best summarizes returns in this particular short window?*
- Once that local factor space is estimated using large N , the paper can recover the local ex post risk premia and SDF and then analytically correct for the bias created by finite T .
- This is why the method is especially valuable when time variation is real but hard to parametrize, and when preselecting instruments is itself risky.

7.3 Takeaway

This paper is best read as a fixed- T , large- N foundation for conditional factor-model inference. Its main message is that time variation in asset-pricing models should often be handled locally rather than by imposing a rigid full-sample parametric structure. The payoff is both econometric and

economic: the procedure is robust to omitted factors and misspecified conditioning information, and it delivers visibly better pricing performance in simulations and in US stock-return data. For research on conditional factor models, latent risk extraction, asset-pricing tests, or rolling-window portfolio construction, this paper provides a clean and powerful framework.